

---

# Advanced Prediction Models

Deep Learning, Graphical Models and Reinforcement  
Learning

# Today's Outline

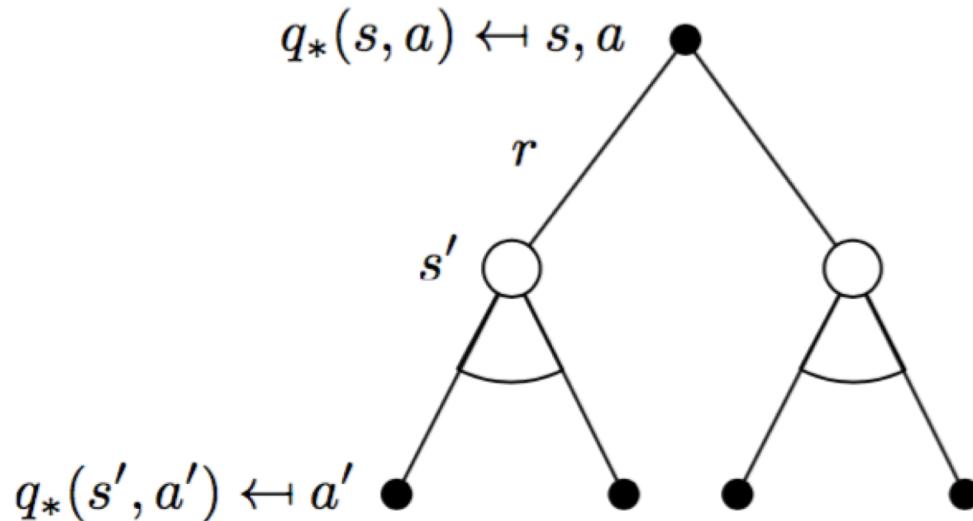
---

- Value Function Approximation
- Deep Reinforcement Learning
  - DQN for Atari Games
  - AlphaGo for Go

# Value Function Approximation

# The Q Learning Algorithm

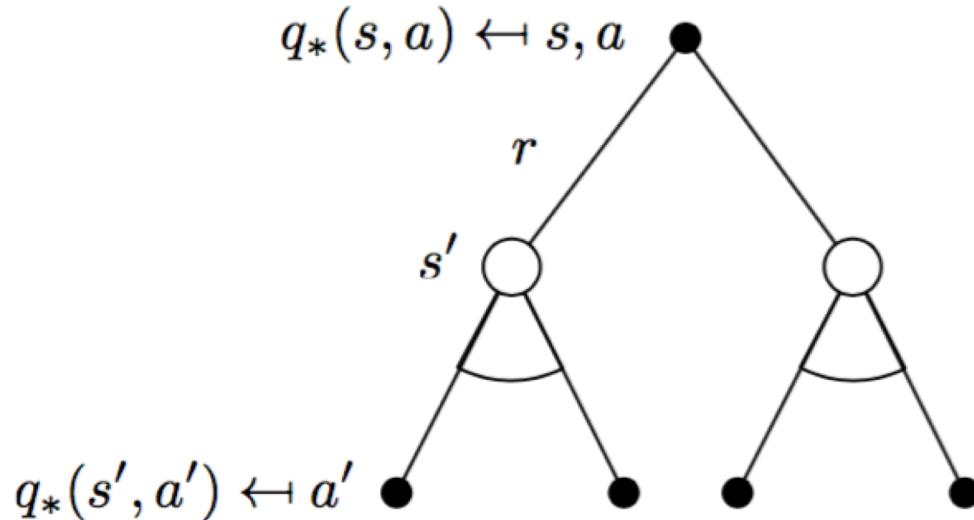
- If we know the model
  - Turn the Bellman Optimality Equation into an **iterative update**
  - This is called Value Iteration



$$q_*(s, a) = \boxed{\mathcal{R}_s^a} + \gamma \sum_{s' \in S} \mathcal{P}_{ss'}^a \max_{a'} q_*(s', a')$$

# The Q Learning Algorithm

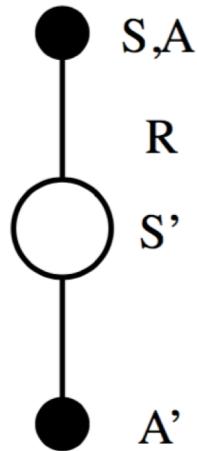
- If we do not know the model
  - Do **sampling** to get an **incremental** iterative update
  - Choose next actions to ensure **exploration**



$$q_*(s, a) = \mathcal{R}_s^a + \gamma \sum_{s' \in S} \mathcal{P}_{ss'}^a \max_{a'} q_*(s', a')$$

# The Q Learning Algorithm

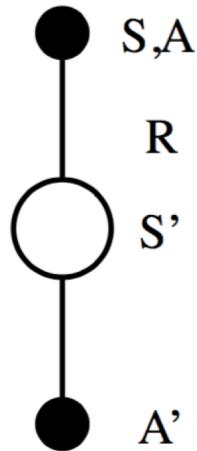
- If we do not know the model
  - Do **sampling** to get an **incremental** iterative update
  - Choose next actions to ensure **exploration**



$$Q(S, A) \leftarrow Q(S, A) + \alpha (R + \gamma Q(S', A') - Q(S, A))$$

# The Q Learning Algorithm

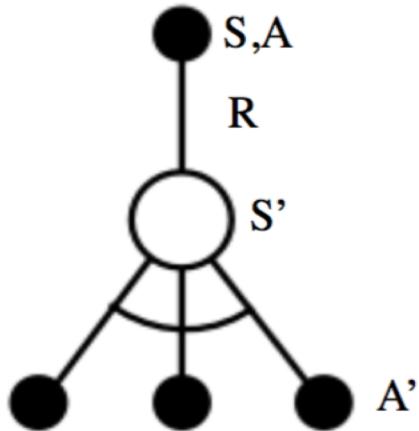
- If we do not know the model
  - Do **sampling** to get an **incremental** iterative update
  - Choose next actions to ensure **exploration**



$$Q(S, A) \leftarrow Q(S, A) + \alpha (R + \gamma Q(S', A') - Q(S, A))$$

# The Q Learning Algorithm

- If we do not know the model
  - Do **sampling** to get an **incremental** iterative update
  - Choose next actions to ensure **exploration**



$$Q(S, A) \leftarrow Q(S, A) + \alpha \left( R + \gamma \max_{a'} Q(S', a') - Q(S, A) \right)$$

# The Q Learning Algorithm

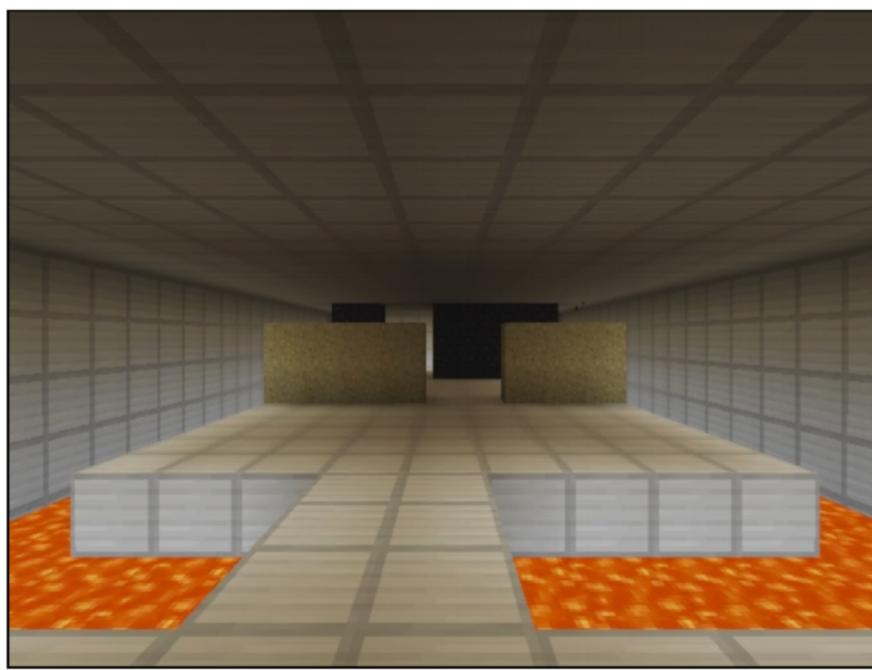
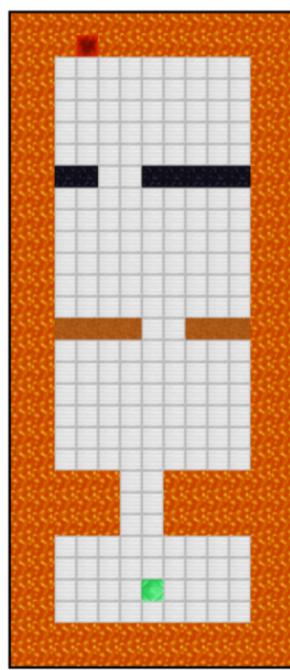
- Initialize  $Q$ , which is a **table** of size #states × #actions
- Start at state  $S_1$
- For  $t = 1, 2, 3, \dots$ 
  - Take  $A_t$  chosen uniformly at random with probability  $\epsilon$
  - Take  $\text{argmax}_{a \in A} Q(S_t, a)$  with probability  $1 - \epsilon$
  - Update  $Q$ :
    - $$Q(S_t, A_t) = Q(S_t, A_t) + \alpha_t (R_{t+1} + \gamma \max_{a \in A} Q(S_{t+1}, a) - Q(S_t, A_t))$$

Explore  
Exploit

Temporal difference error
- Parameter  $\epsilon$  is the exploration parameter
- Parameter  $\alpha_t$  is the learning rate
- Under appropriate assumptions<sup>1</sup>,  $\lim_{t \rightarrow \infty} Q = Q^*$

<sup>1</sup>Reference: Christopher J. C. H. Watkins and Peter Dayan, 1992

# Tabular Q Learning is Not Enough



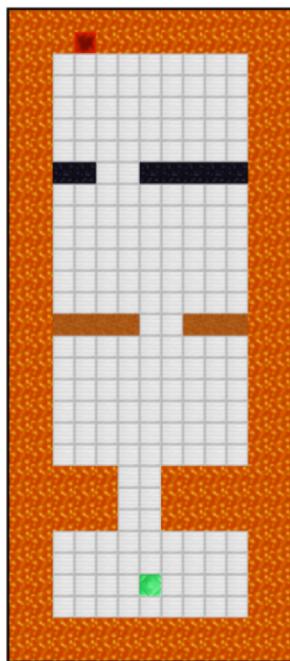
Robotic agent navigating in real-world (left)

States: Position in a grid

Actions: Forward/Back/Left/Right

Reward: 1 on reaching target, -100 for dying

# Tabular Q Learning is Not Enough



Robotic agent navigating in real-world (right)

States: Camera view in front of the robot

Actions: Forward/Back/Left/Right

Reward: 1 on reaching target, -100 for dying

# Function Approximation Recipe

---

- Use a deep network or any other function class to represent
  - the value function, and/or
  - the policy, and/or
  - the model

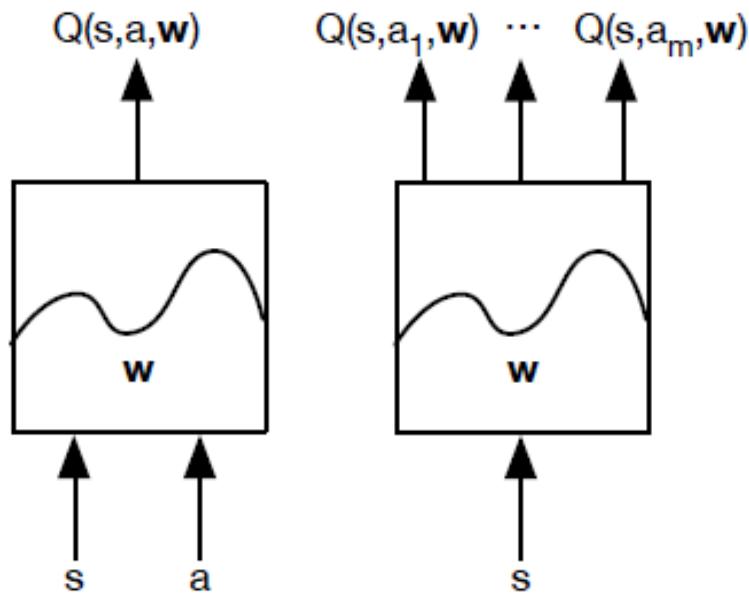
# Function Approximation Recipe

---

- Use a deep network or any other function class to represent
  - the value function, and/or
  - the policy, and/or
  - the model
- Optimize this network end to end
  - Example:
    - If the approximator is differentiable
    - Use stochastic gradient descent
- Do the optimization incrementally or in batch mode

# Q Function Approximation

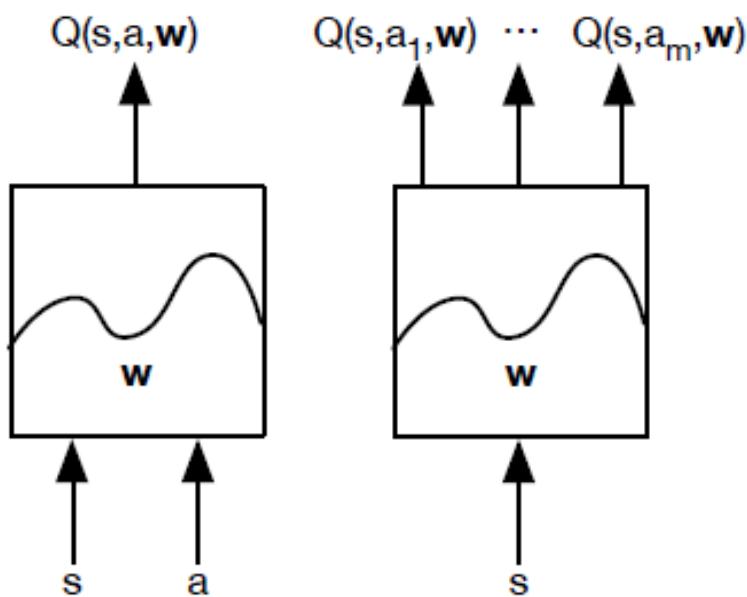
- Instead of storing  $\# \text{states} \times \# \text{action parameters}$  in a table, we want to find more scalable ways to capture Q values
- Represent  $Q$  using a function approximator with weights  $w$ :  
$$Q(s, a; w) \approx Q^*(s, a)$$



<sup>1</sup>Figure: David Silver

# Q Function Approximation

- Instead of storing  $\# \text{states} \times \# \text{action parameters}$  in a table, we want to find more scalable ways to capture Q values
- Represent  $Q$  using a function approximator with weights  $w$ :  
$$Q(s, a; w) \approx Q^*(s, a)$$



Linear

Decision tree

Neural network

Basis functions

Nearest neighbor

# Q Function Approximation

- Approximate the action-value function

$$\hat{q}(S, A, \mathbf{w}) \approx q_{\pi}(S, A)$$

# Q Function Approximation

- Approximate the action-value function

$$\hat{q}(S, A, \mathbf{w}) \approx q_\pi(S, A)$$

- Minimise mean-squared error between approximate action-value fn  $\hat{q}(S, A, \mathbf{w})$  and true action-value fn  $q_\pi(S, A)$

$$J(\mathbf{w}) = \mathbb{E}_\pi [(q_\pi(S, A) - \hat{q}(S, A, \mathbf{w}))^2]$$

# Q Function Approximation

- Approximate the action-value function

$$\hat{q}(S, A, \mathbf{w}) \approx q_\pi(S, A)$$

- Minimise mean-squared error between approximate action-value fn  $\hat{q}(S, A, \mathbf{w})$  and true action-value fn  $q_\pi(S, A)$

$$J(\mathbf{w}) = \mathbb{E}_\pi [(q_\pi(S, A) - \hat{q}(S, A, \mathbf{w}))^2]$$

- Use stochastic gradient descent to find a local minimum

$$-\frac{1}{2} \nabla_{\mathbf{w}} J(\mathbf{w}) = (q_\pi(S, A) - \hat{q}(S, A, \mathbf{w})) \nabla_{\mathbf{w}} \hat{q}(S, A, \mathbf{w})$$

$$\Delta \mathbf{w} = \alpha (q_\pi(S, A) - \hat{q}(S, A, \mathbf{w})) \nabla_{\mathbf{w}} \hat{q}(S, A, \mathbf{w})$$

# Q Function Approximation: Example

- Represent state *and* action by a *feature vector*

$$\mathbf{x}(S, A) = \begin{pmatrix} \mathbf{x}_1(S, A) \\ \vdots \\ \mathbf{x}_n(S, A) \end{pmatrix}$$

# Q Function Approximation: Example

- Represent state *and* action by a *feature vector*

$$\mathbf{x}(S, A) = \begin{pmatrix} \mathbf{x}_1(S, A) \\ \vdots \\ \mathbf{x}_n(S, A) \end{pmatrix}$$

- Represent action-value fn by linear combination of features

$$\hat{q}(S, A, \mathbf{w}) = \mathbf{x}(S, A)^\top \mathbf{w} = \sum_{j=1}^n \mathbf{x}_j(S, A) \mathbf{w}_j$$

# Q Function Approximation: Example

- Represent state *and* action by a *feature vector*

$$\mathbf{x}(S, A) = \begin{pmatrix} \mathbf{x}_1(S, A) \\ \vdots \\ \mathbf{x}_n(S, A) \end{pmatrix}$$

- Represent action-value fn by linear combination of features

$$\hat{q}(S, A, \mathbf{w}) = \mathbf{x}(S, A)^\top \mathbf{w} = \sum_{j=1}^n \mathbf{x}_j(S, A) w_j$$

- Stochastic gradient descent update

$$\nabla_{\mathbf{w}} \hat{q}(S, A, \mathbf{w}) = \mathbf{x}(S, A)$$

$$\Delta \mathbf{w} = \alpha (q_\pi(S, A) - \hat{q}(S, A, \mathbf{w})) \mathbf{x}(S, A)$$

# Q Function Approximation: Another Perspective

---

- Recall the Q Learning update

$$Q(S_t, A_t) = Q(S_t, A_t) + \alpha_t(R_{t+1} + \gamma \max_{a \in A} Q(S_{t+1}, a) - Q(S_t, A_t))$$

# Q Function Approximation: Another Perspective

- Recall the Q Learning update

$$Q(S_t, A_t) = Q(S_t, A_t) + \alpha_t (R_{t+1} + \gamma \max_{a \in A} Q(S_{t+1}, a) - Q(S_t, A_t))$$

- At optimality

- $E \left[ R_{t+1} + \gamma \max_{a \in A} Q(S_{t+1}, a) - Q(S_t, A_t) \right] = 0$

# Q Function Approximation: Another Perspective

- Recall the Q Learning update

$$Q(S_t, A_t) = Q(S_t, A_t) + \alpha_t (R_{t+1} + \gamma \max_{a \in A} Q(S_{t+1}, a) - Q(S_t, A_t))$$

- At optimality
  - $E \left[ R_{t+1} + \gamma \max_{a \in A} Q(S_{t+1}, a) - Q(S_t, A_t) \right] = 0$
- Intuitively, this tells us to minimize the empirical error between
  - $R_{t+1} + \gamma \max_{a \in A} Q(S_{t+1}, a, w)$  and  $Q(S_t, A_t, w)$

# Example: Function Approximation Success (2013)



<sup>1</sup>Figure: Defazio Graepel, Atari Learning Environment

# Issues with Function Approximation

---

- This can potentially be a nonlinear optimization over  $W$ 
  - Unless we use a linear approximator
- Can optimize incrementally or in batch
  - Which is better? (we will answer *this* for DQN later)
- Naïve optimization may diverge and oscillate! This is because
  - The data is not i.i.d.
  - Policy/Value may be too sensitive to action choice (max over actions may completely change future trajectory)

---

---

# Questions?

# Today's Outline

---

- Value Function Approximation
- Deep Reinforcement Learning
  - DQN for Atari Games
  - AlphaGo for Go

---

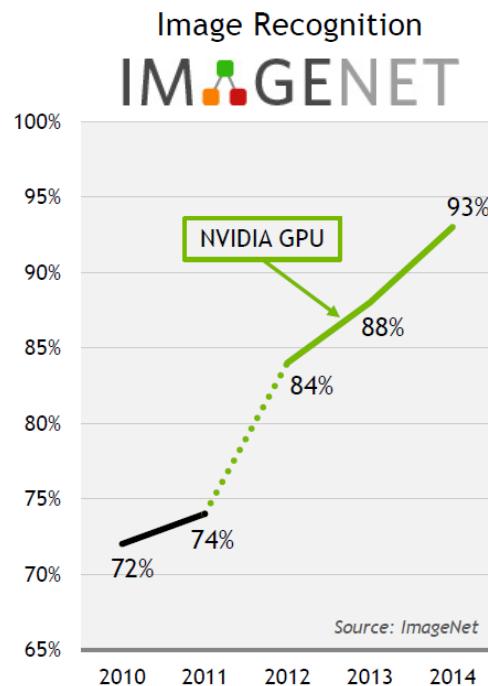
# Deep Reinforcement Learning I: DQN

# Deep Reinforcement Learning

---

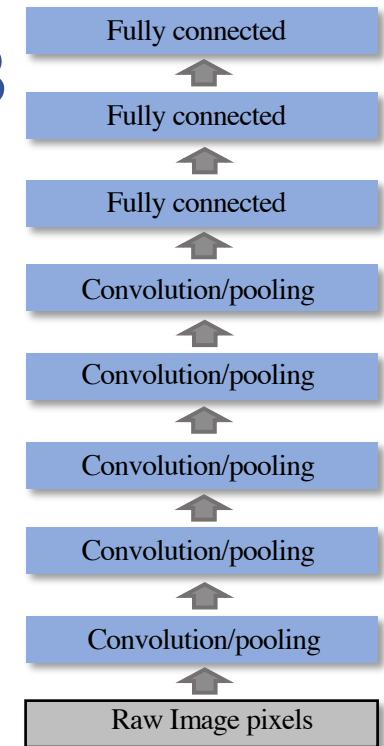
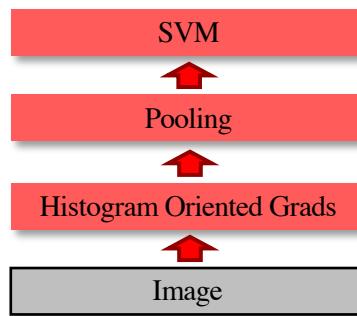
- Bringing in the success of deep perception/prediction architectures to function approximation
- We will look at two RL agents
  - DQN (2013)
  - AlphaGo (2016)
- Attempt to highlight some additional aspects that made these agents succeed so well in their respective domains

# Why Deep Representations?



2012-2013

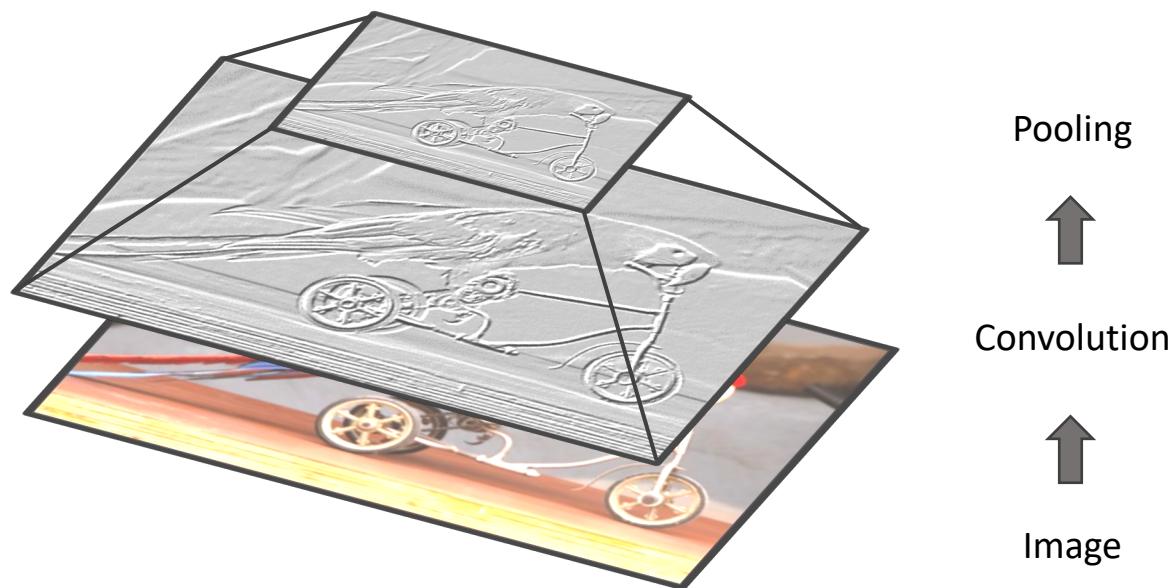
earlier



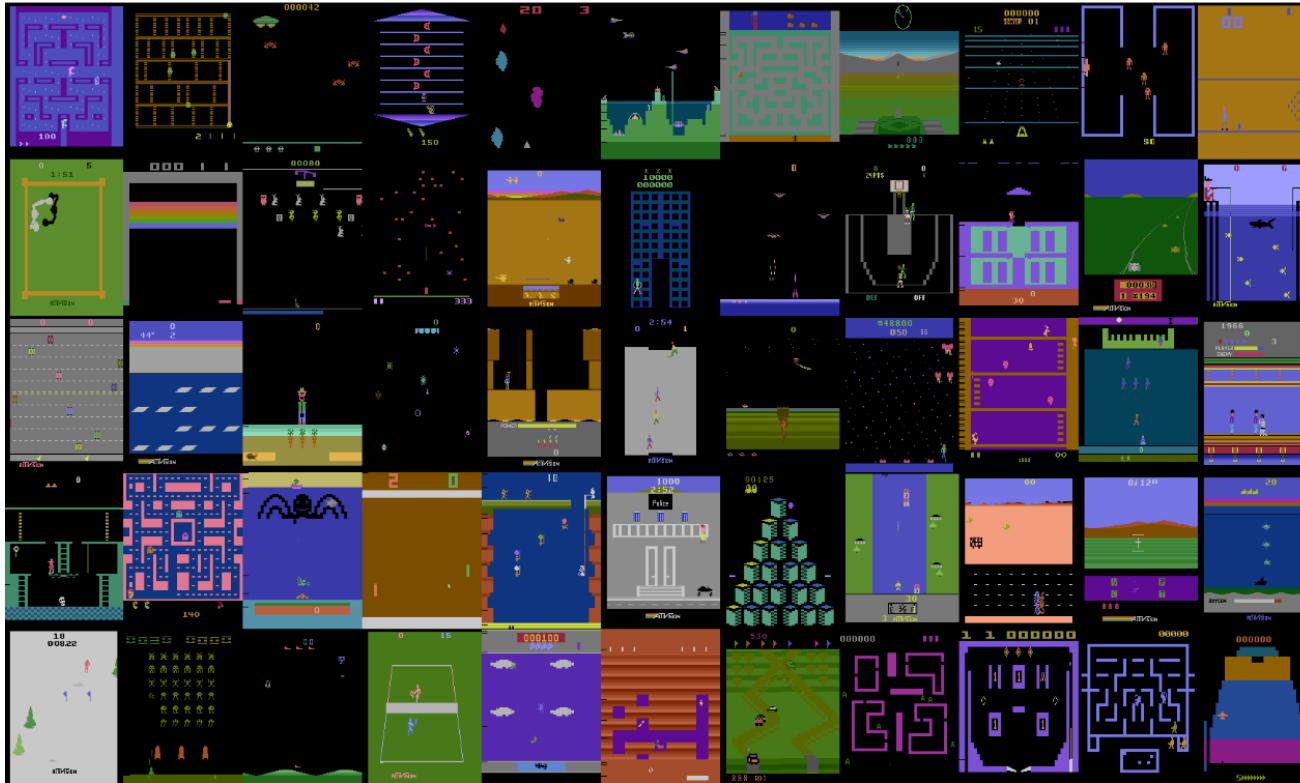
<sup>1</sup>Reference: Julie Bernauer/Ryan Olson, Li Deng

# Why Deep Representations?

- CNN as the Function Approximator
- Captures two key properties
  - Local connections with weight sharing
  - Pooling for translation invariance

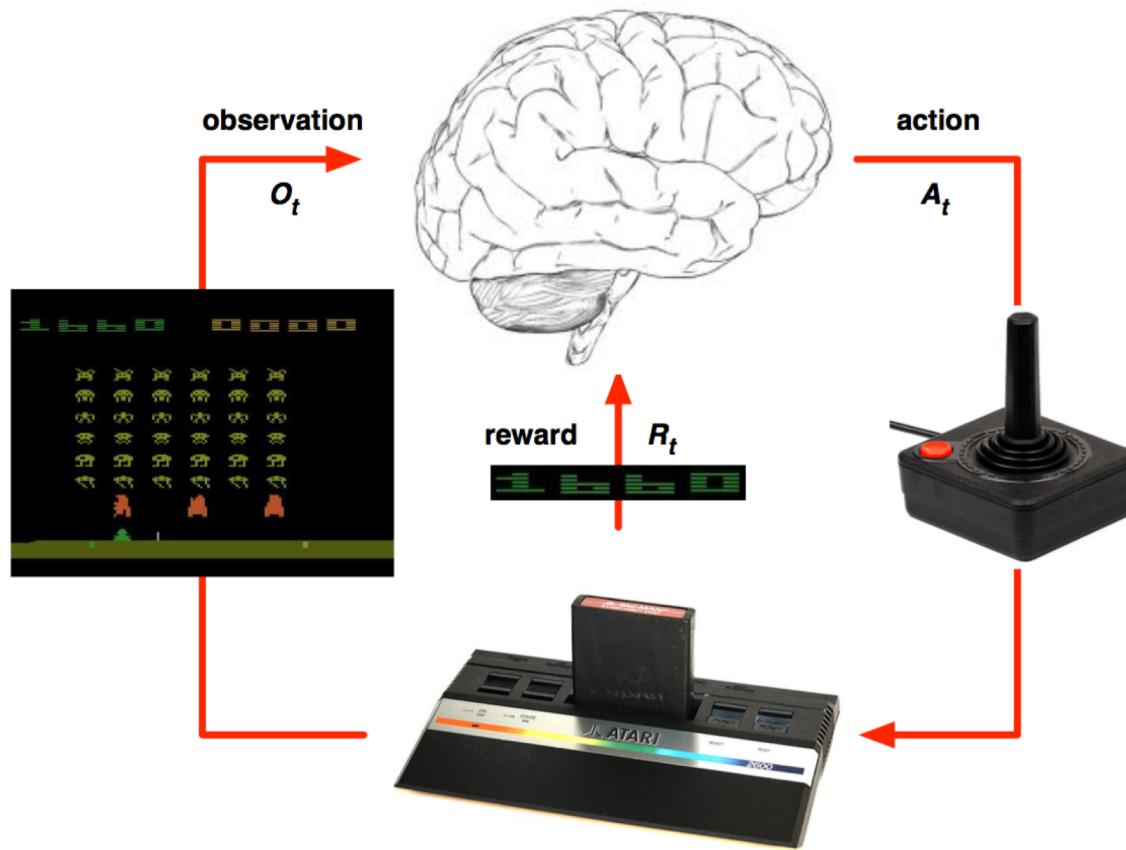


# DQN Plays Atari (2013)



<sup>1</sup>Figure: Defazio Graepel, Atari Learning Environment

# DQN Architecture



- Rules of the game are unknown
- Learn directly from interactive game-play
- Pick actions on joystick, see pixels and scores

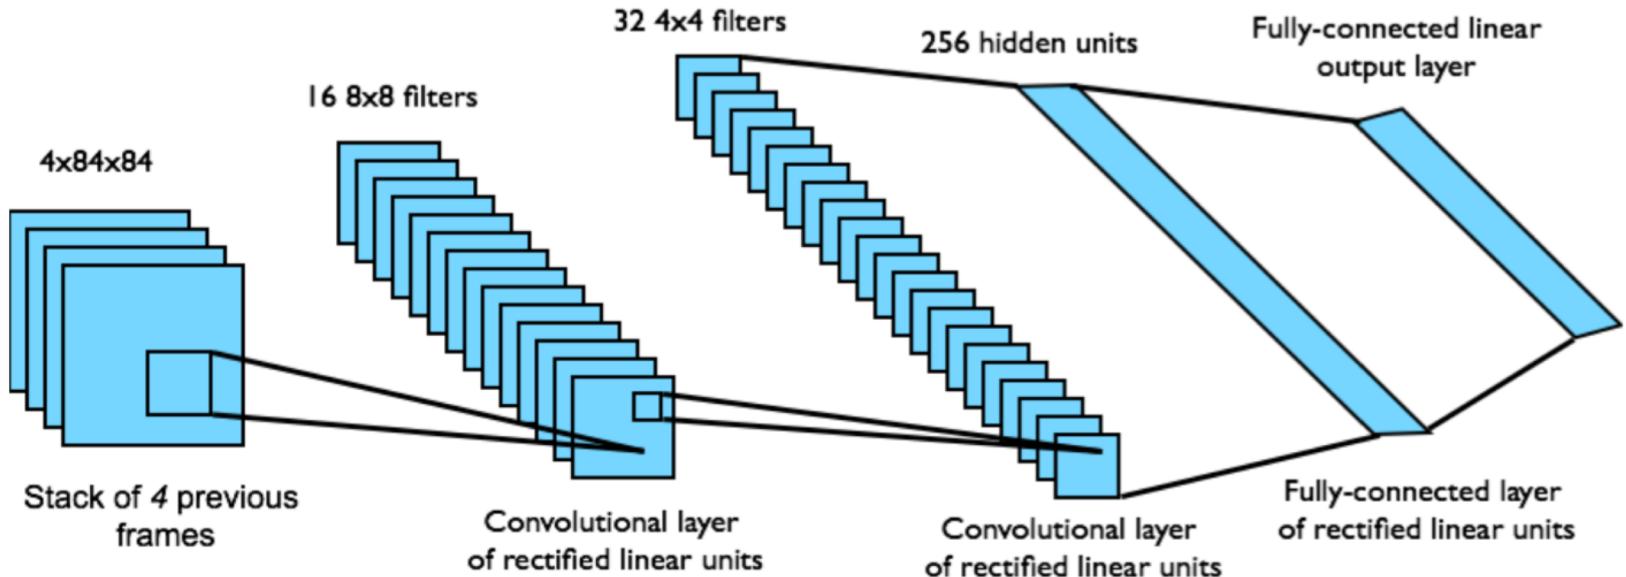
# DQN Extends Function Approximation

---

- DQN does Q learning with function approximation
- Uses a CNN as the approximator
- Extension
  - Does batch optimization to update the weights
  - Freezes targets over several steps

# DQN Extends Function Approximation

- End-to-end learning of values  $Q(s, a)$  from pixels  $s$
- Input state  $s$  is stack of raw pixels from last 4 frames
- Output is  $Q(s, a)$  for 18 joystick/button positions
- Reward is change in score for that step



# DQN Extends Function Approximation

---

DQN uses **experience replay** and **fixed Q-targets**

- Take action  $a_t$  according to  $\epsilon$ -greedy policy

# DQN Extends Function Approximation

DQN uses **experience replay** and **fixed Q-targets**

- Take action  $a_t$  according to  $\epsilon$ -greedy policy
- Store transition  $(s_t, a_t, r_{t+1}, s_{t+1})$  in replay memory  $\mathcal{D}$
- Sample random mini-batch of transitions  $(s, a, r, s')$  from  $\mathcal{D}$

# DQN Extends Function Approximation

DQN uses **experience replay** and **fixed Q-targets**

- Take action  $a_t$  according to  $\epsilon$ -greedy policy
- Store transition  $(s_t, a_t, r_{t+1}, s_{t+1})$  in replay memory  $\mathcal{D}$
- Sample random mini-batch of transitions  $(s, a, r, s')$  from  $\mathcal{D}$
- Compute Q-learning targets w.r.t. old, fixed parameters  $w^-$

# DQN Extends Function Approximation

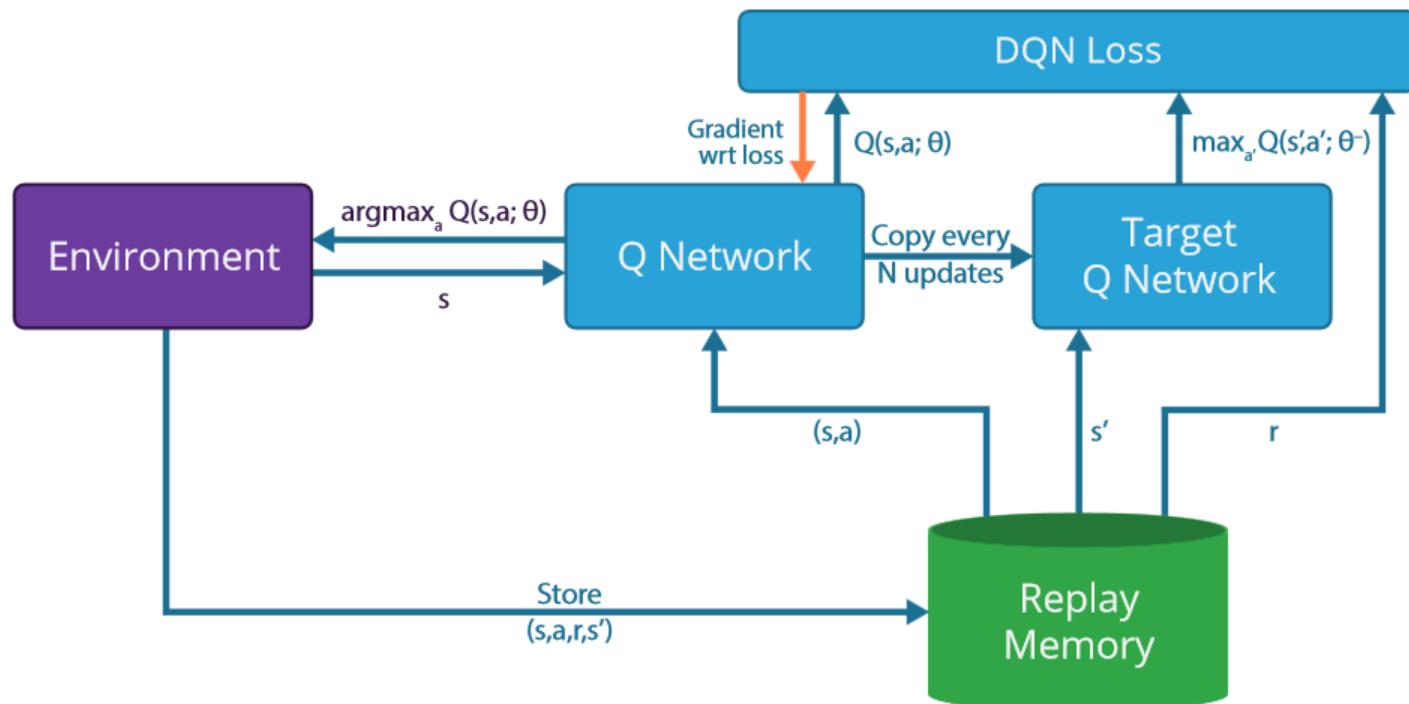
DQN uses **experience replay** and **fixed Q-targets**

- Take action  $a_t$  according to  $\epsilon$ -greedy policy
- Store transition  $(s_t, a_t, r_{t+1}, s_{t+1})$  in replay memory  $\mathcal{D}$
- Sample random mini-batch of transitions  $(s, a, r, s')$  from  $\mathcal{D}$
- Compute Q-learning targets w.r.t. old, fixed parameters  $w^-$
- Optimise MSE between Q-network and Q-learning targets

$$\mathcal{L}_i(w_i) = \mathbb{E}_{s,a,r,s' \sim \mathcal{D}_i} \left[ \left( r + \gamma \max_{a'} Q(s', a'; w_i^-) - Q(s, a; w_i) \right)^2 \right]$$

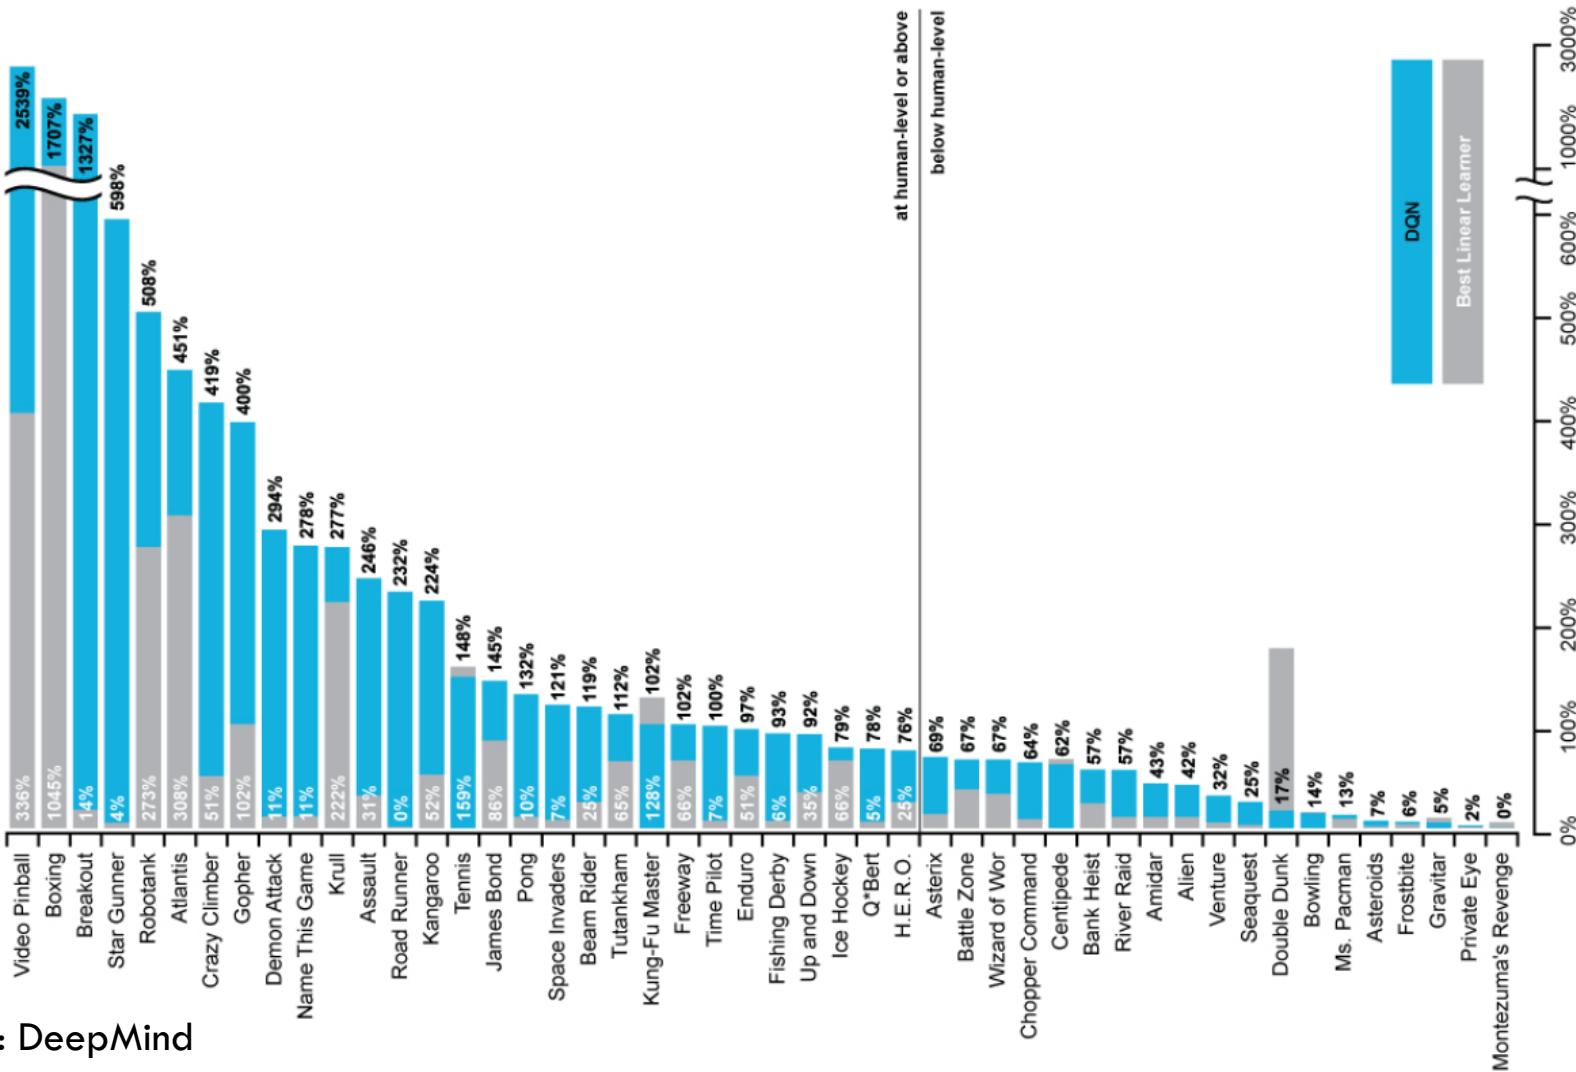
- Using variant of stochastic gradient descent

# DQN Architecture



# DQN Performance Results

- DQN does not know the rules of the game a-priori
- No feature engineering or hyper-parameter tuning for DQN across games

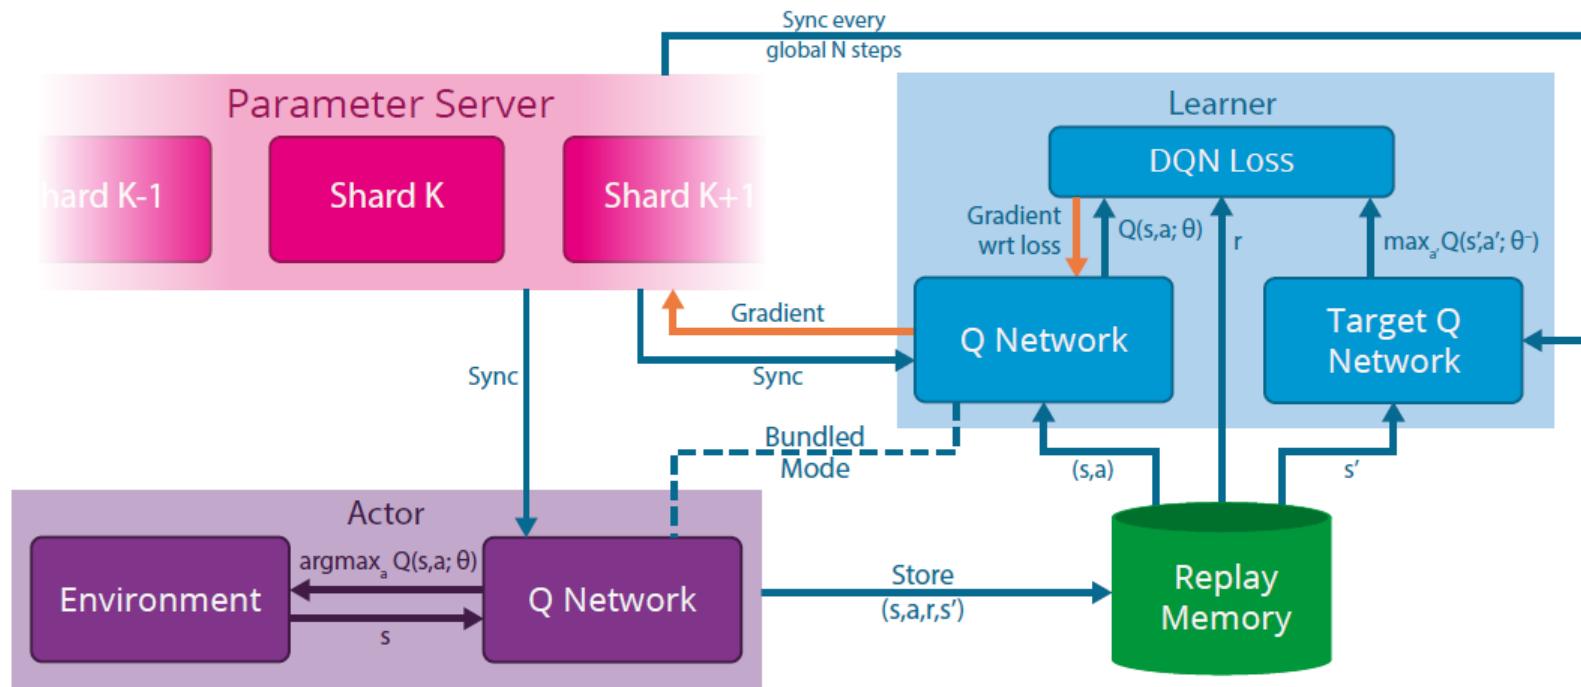


<sup>1</sup>Figure: DeepMind

# Did the Extensions Help?

	Replay Fixed-Q	Replay Q-learning	No replay Fixed-Q	No replay Q-learning
Breakout	316.81	240.73	10.16	3.17
Enduro	1006.3	831.25	141.89	29.1
River Raid	7446.62	4102.81	2867.66	1453.02
Seaquest	2894.4	822.55	1003	275.81
Space Invaders	1088.94	826.33	373.22	301.99

# Scalable Version: An Architecture by Google



- 100 actors, 100 learners, and 31 parameter holding machines.
- Reduce compute from 14 days to 6 hours
- This is a 30x speedup using 200x compute power

---

---

# Questions?

# Today's Outline

---

- Value Function Approximation
- Deep Reinforcement Learning
  - DQN for Atari Games
  - AlphaGo for Go

---

# Deep Reinforcement Learning I: AlphaGo

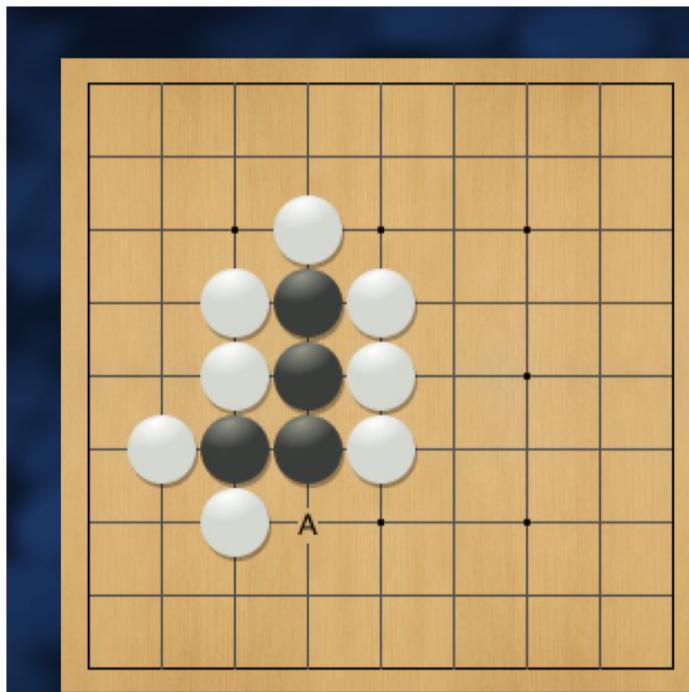
# AlphaGo Conquers Go (2016)



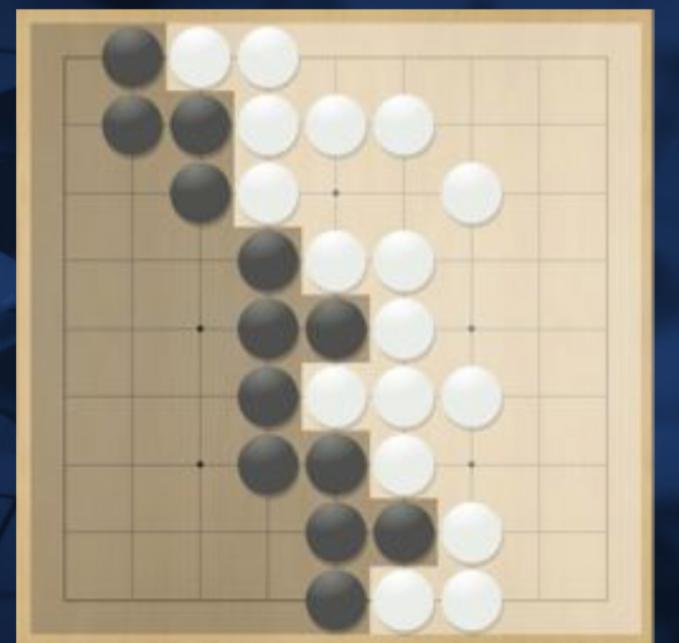
<sup>1</sup>Reference: DeepMind, March 2016

# The Game of Go

- Go is 2500 years old. Has about  $10^{270}$  states.
- Making it impossible for computers to evaluate who is winning



Capture



Territory

# The Game of Go

- Go was one of the only classic board games before March 2016, where AI agents were not the best

Program	Level of Play	RL Program to Achieve Level
Checkers	Perfect	<i>Chinook</i>
Chess	International Master	<i>KnightCap / Meep</i>
Othello	Superhuman	<i>Logistello</i>
Backgammon	Superhuman	<i>TD-Gammon</i>
Scrabble	Superhuman	<i>Maven</i>
Go	Grandmaster	<i>MoGo<sup>1</sup>, Crazy Stone<sup>2</sup>, Zen<sup>3</sup></i>
Poker <sup>4</sup>	Superhuman	<i>SmooCT</i>

---

<sup>1</sup>9 × 9

<sup>2</sup>9 × 9 and 19 × 19

<sup>3</sup>19 × 19

<sup>4</sup>Heads-up Limit Texas Hold'em

# The Forward Search Problem

---

- Recall the two sequential decision making problems
  - Reinforcement learning
  - Planning
- The forward search problem is a planning problem
  - That is, we know the model of the world
- Useful in the case when we cannot plan everything beforehand
- Focus on what action to take next

# The Forward Search Problem for Go

- How good is a position  $s$ ?
- Reward function (undiscounted):

$R_t = 0$  for all non-terminal steps  $t < T$

$$R_T = \begin{cases} 1 & \text{if Black wins} \\ 0 & \text{if White wins} \end{cases}$$

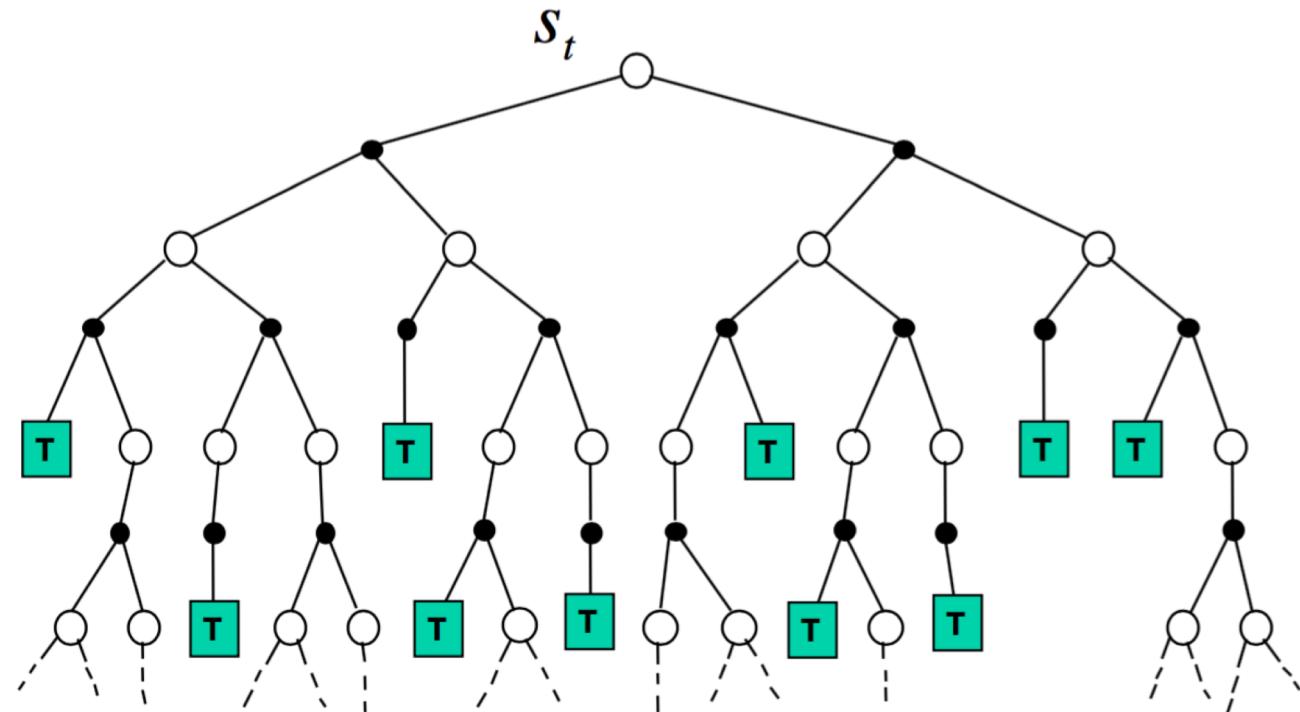
- Policy  $\pi = \langle \pi_B, \pi_W \rangle$  selects moves for both players
- Value function (how good is position  $s$ ):

$$v_\pi(s) = \mathbb{E}_\pi [R_T \mid S = s] = \mathbb{P} [\text{Black wins} \mid S = s]$$

$$v_*(s) = \max_{\pi_B} \min_{\pi_W} v_\pi(s)$$

# Forward Search Using Simulations

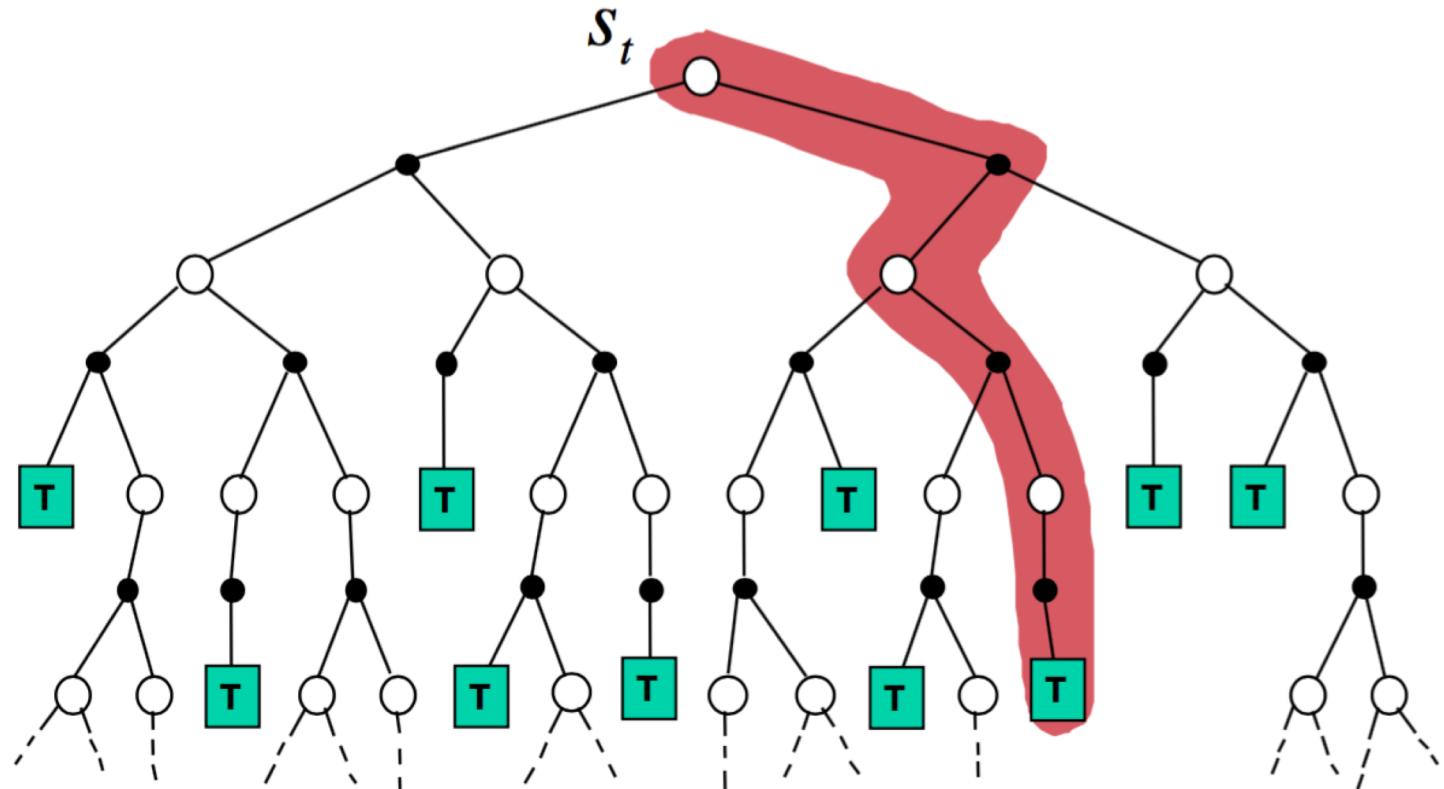
- Forward search algorithms select the best action by lookahead
- They build a search tree with the current state  $s_t$  at the root
- Using a model of the MDP to look ahead



- No need to solve whole MDP, just sub-MDP starting from now

# Forward Search Using Simulations

- Simulate episodes of experience from **now** with the model
- Apply **model-free** RL to simulated episodes



# Forward Search Using Simulations

- Simulate episodes of experience from now with the model

$$\{s_t^k, A_t^k, R_{t+1}^k, \dots, S_T^k\}_{k=1}^K \sim \mathcal{M}_\nu$$

- Apply model-free RL to simulated episodes

- We will look at two variants
  - Simple Monte Carlo Search
  - Monte Carlo Tree Search

# Simple Monte Carlo Search

- Given a model  $\mathcal{M}_\nu$ , and a **simulation policy**  $\pi$
- For each action  $a \in \mathcal{A}$ 
  - Simulate  $K$  episodes from current (real) state  $s_t$

$$\{\textcolor{red}{s_t}, \textcolor{red}{a}, R_{t+1}^k, S_{t+1}^k, A_{t+1}^k, \dots, S_T^k\}_{k=1}^K \sim \mathcal{M}_\nu, \pi$$

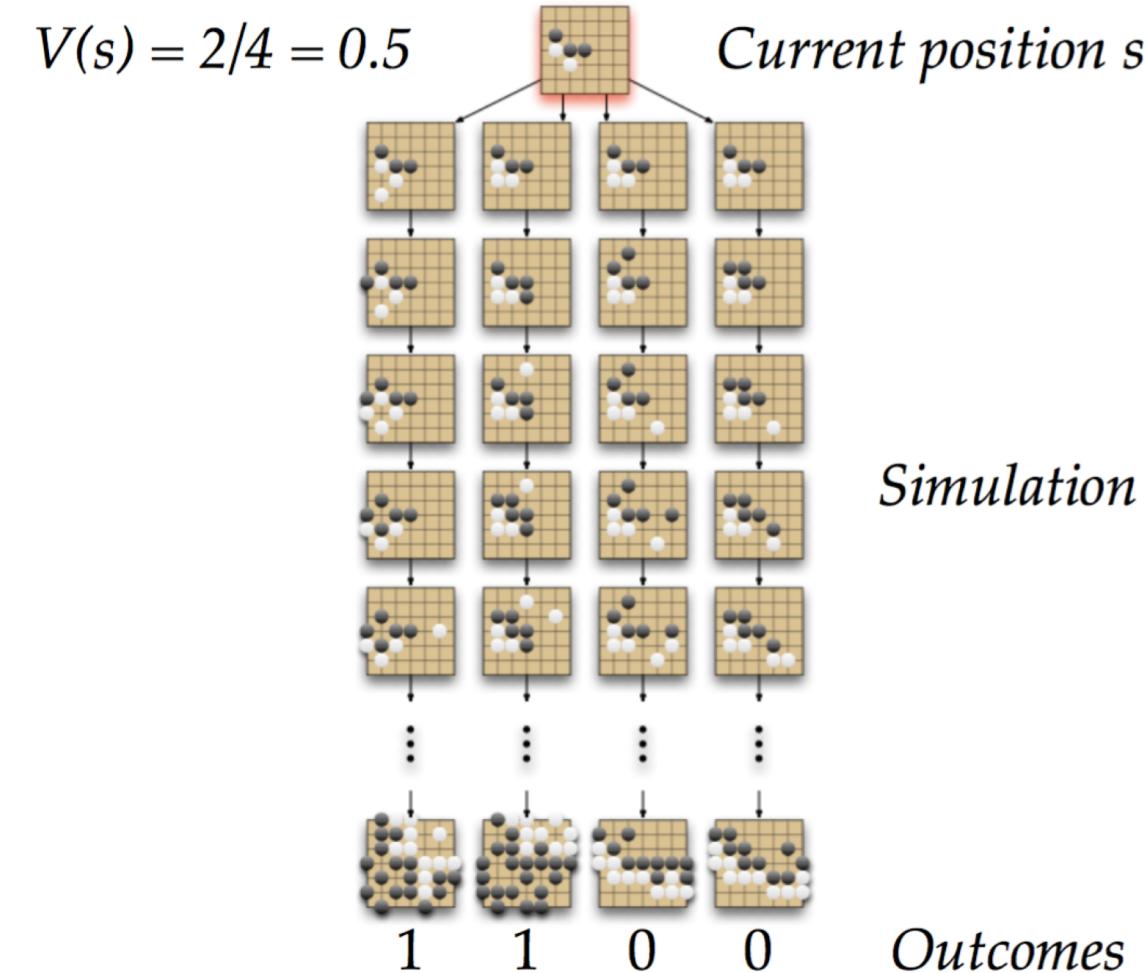
- Evaluate actions by mean return (**Monte-Carlo evaluation**)

$$Q(\textcolor{red}{s_t}, \textcolor{red}{a}) = \frac{1}{K} \sum_{k=1}^K G_t \xrightarrow{P} q_\pi(s_t, a)$$

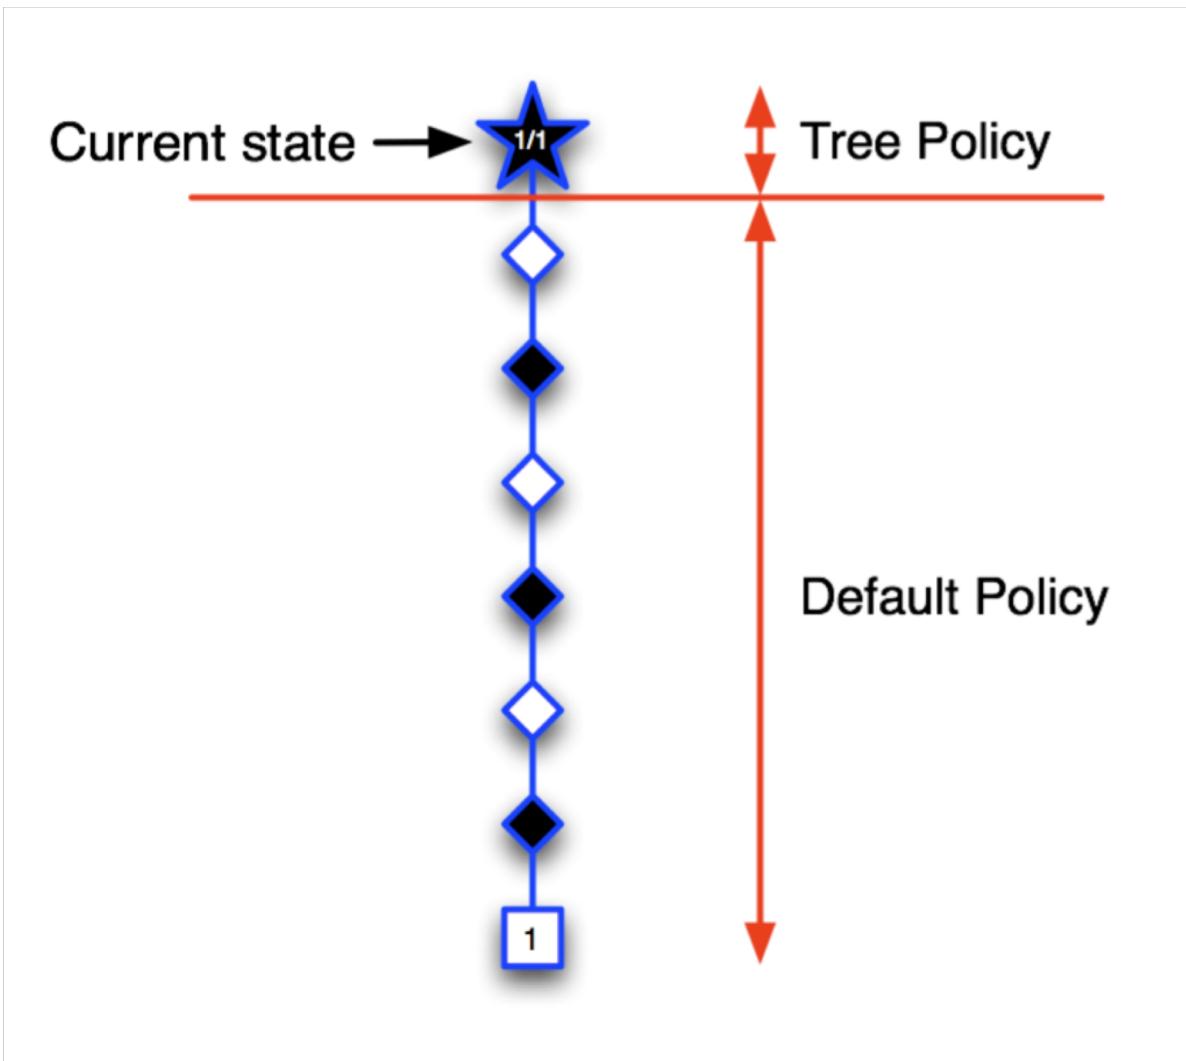
- Select current (real) action with maximum value

$$a_t = \underset{a \in \mathcal{A}}{\operatorname{argmax}} Q(s_t, a)$$

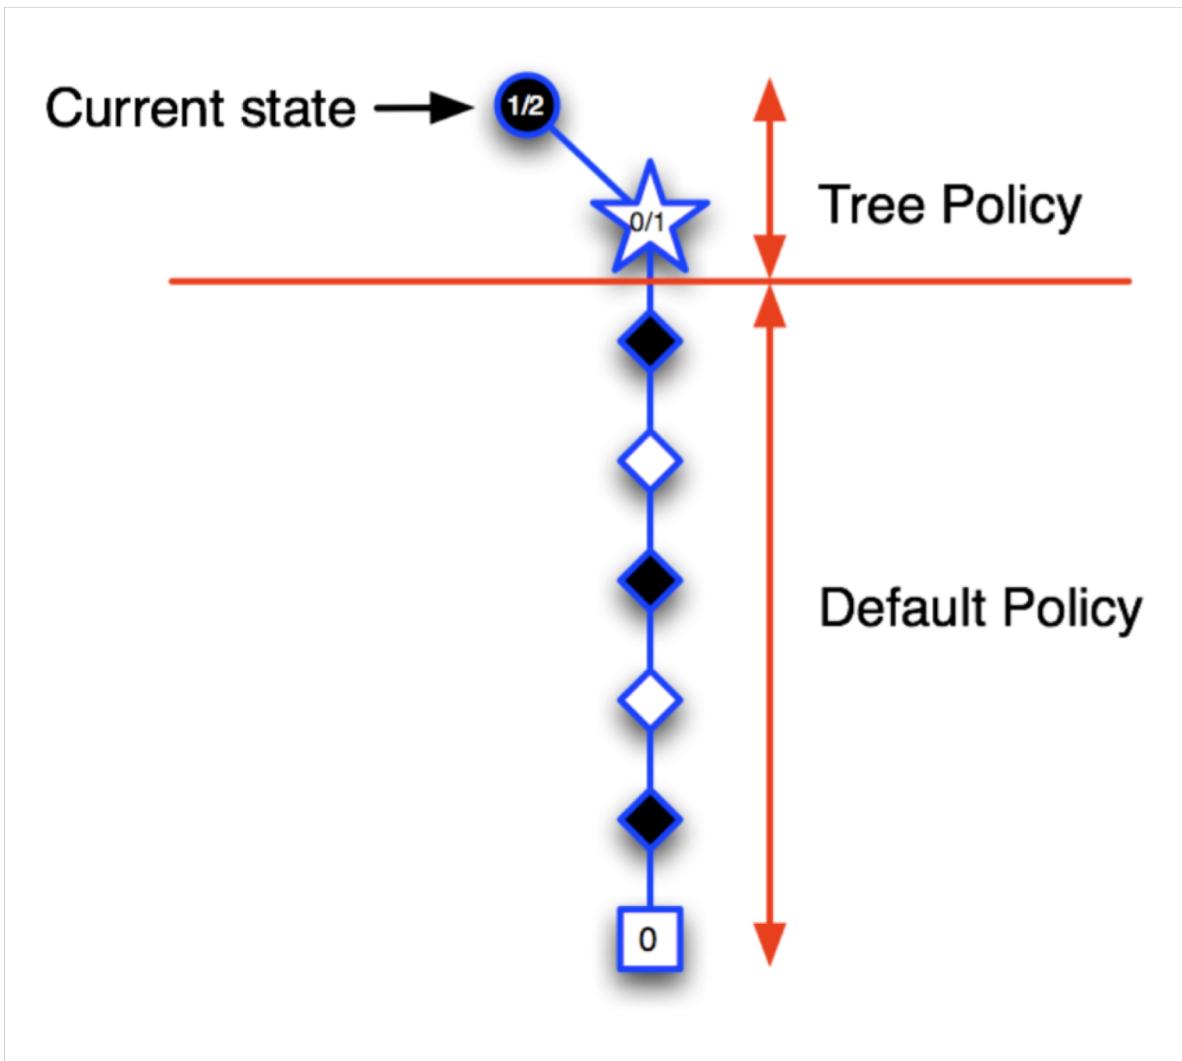
# Simple Monte Carlo Search for Go



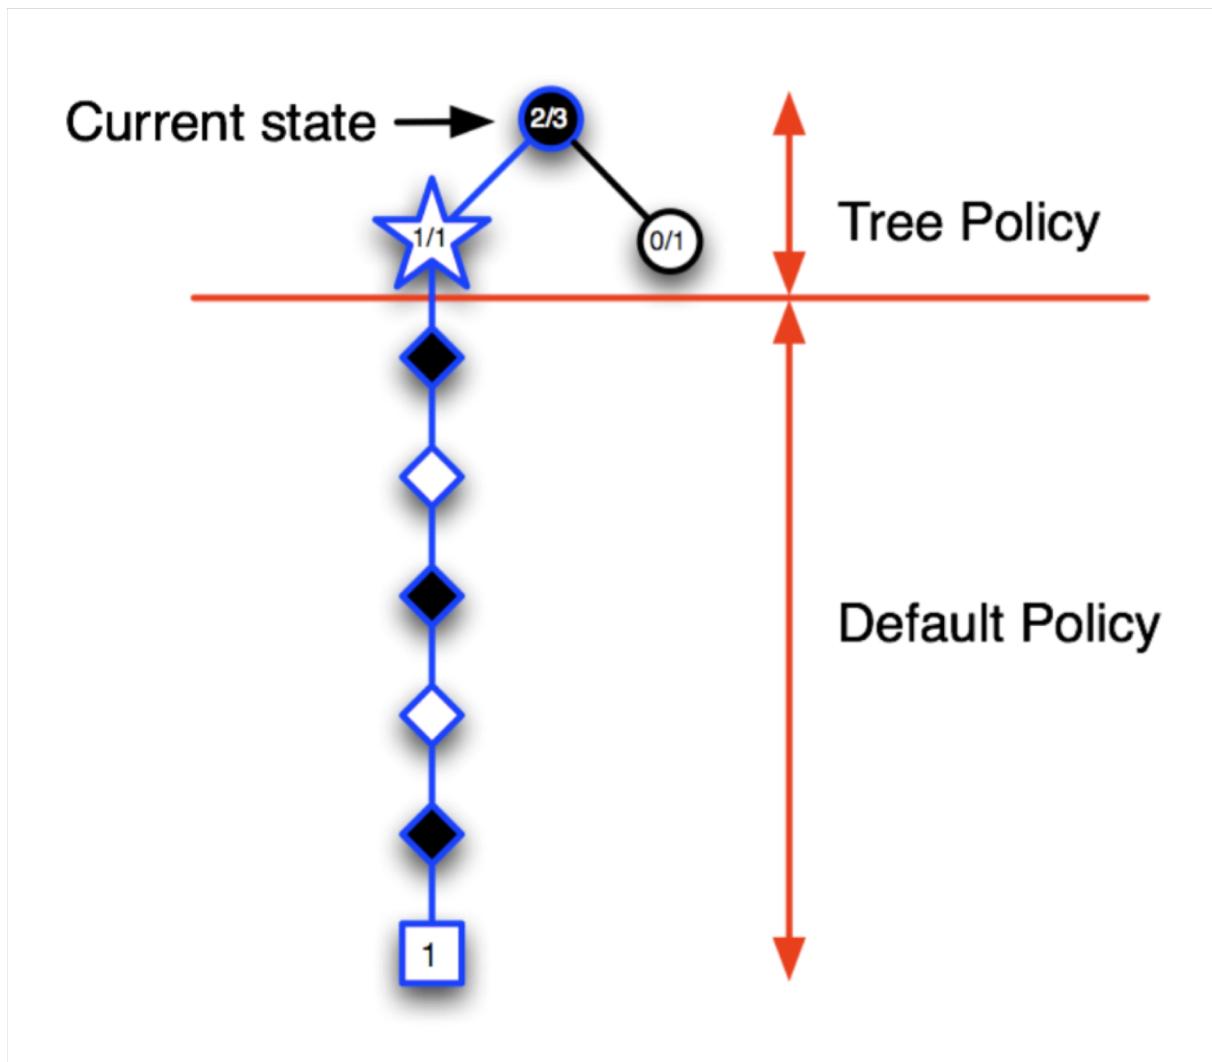
# Monte Carlo Tree Search for Go



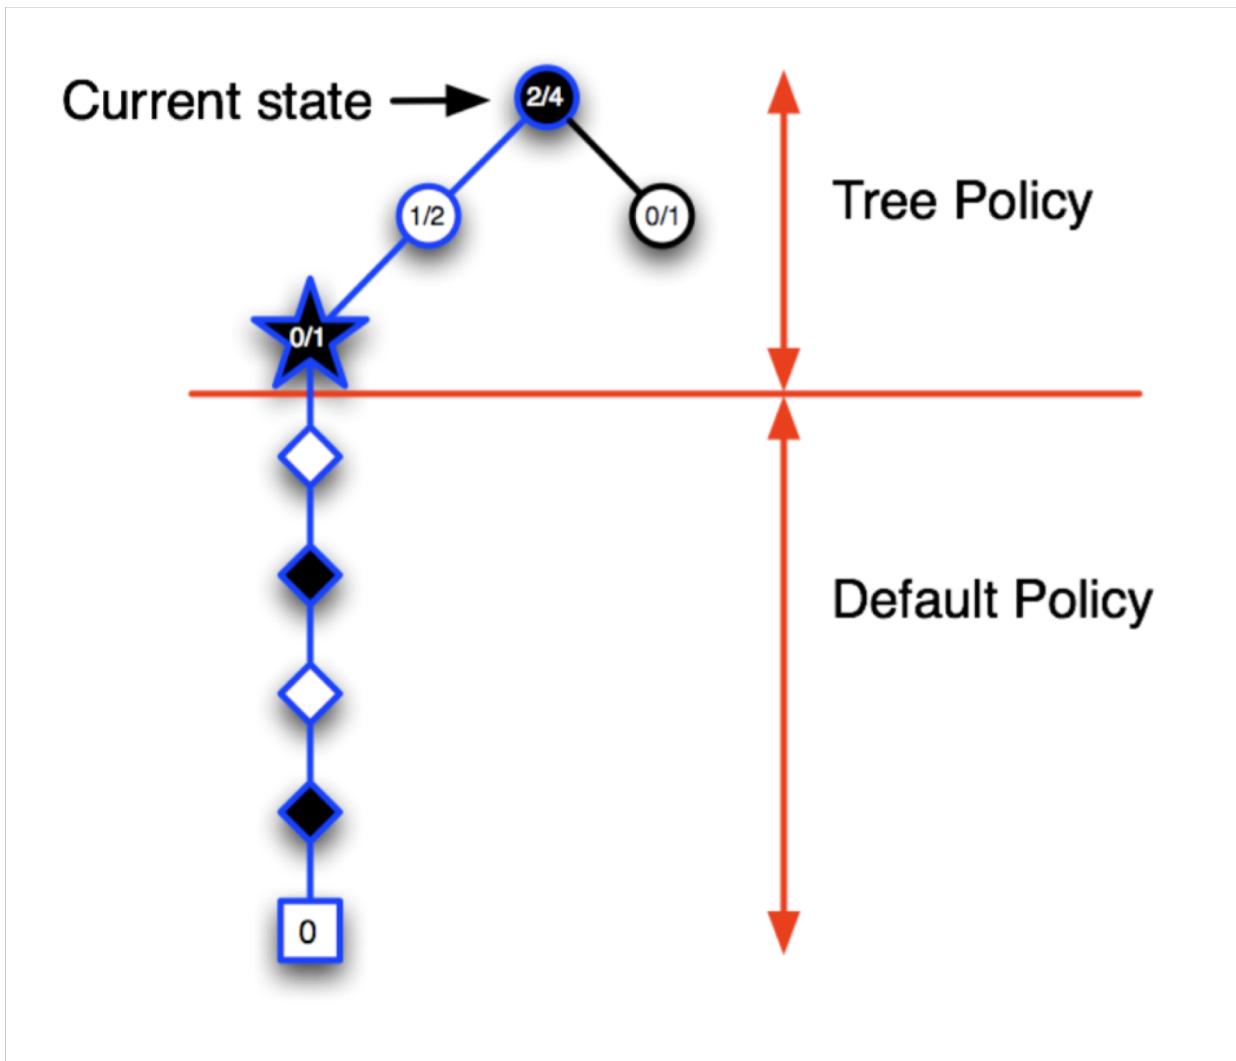
# Monte Carlo Tree Search for Go



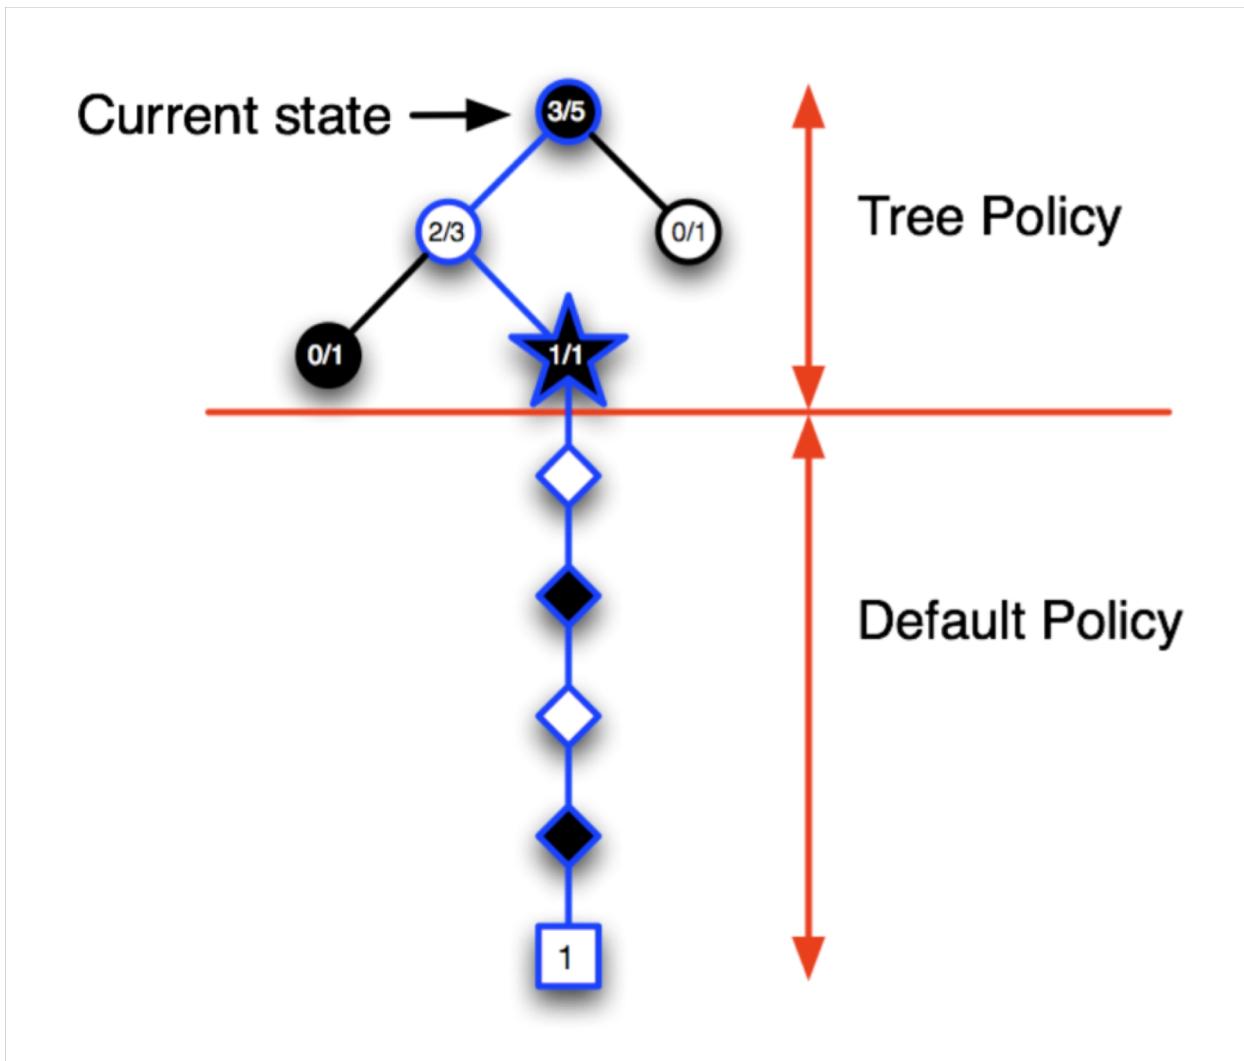
# Monte Carlo Tree Search for Go



# Monte Carlo Tree Search for Go



# Monte Carlo Tree Search for Go



# Monte Carlo Tree Search: Evaluation

- Given a model  $\mathcal{M}_\nu$ ,
- Simulate  $K$  episodes from current state  $s_t$  using current simulation policy  $\pi$

$$\{\textcolor{red}{s_t}, A_t^k, R_{t+1}^k, S_{t+1}^k, \dots, S_T^k\}_{k=1}^K \sim \mathcal{M}_\nu, \pi$$

- Build a search tree containing visited states and actions
- Evaluate states  $Q(s, a)$  by mean return of episodes from  $s, a$

$$Q(\textcolor{red}{s, a}) = \frac{1}{N(s, a)} \sum_{k=1}^K \sum_{u=t}^T \mathbf{1}(S_u, A_u = s, a) G_u \xrightarrow{P} q_\pi(s, a)$$

- After search is finished, select current (real) action with maximum value in search tree

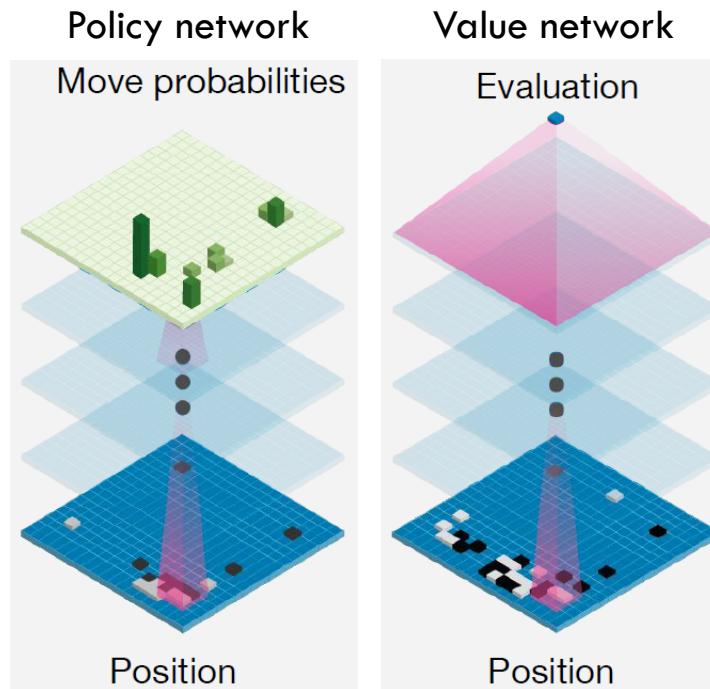
$$a_t = \operatorname{argmax}_{a \in \mathcal{A}} Q(s_t, a)$$

# Monte Carlo Tree Search: Simulation

- In MCTS, the simulation policy  $\pi$  improves
- Each simulation consists of two phases (in-tree, out-of-tree)
  - Tree policy (improves): pick actions to maximise  $Q(S, A)$
  - Default policy (fixed): pick actions randomly
- Repeat (each simulation)
  - Evaluate states  $Q(S, A)$  by Monte-Carlo evaluation
  - Improve tree policy, e.g. by  $\epsilon$  – greedy( $Q$ )
- Monte-Carlo control applied to simulated experience
- Converges on the optimal search tree,  $Q(S, A) \rightarrow q_*(S, A)$

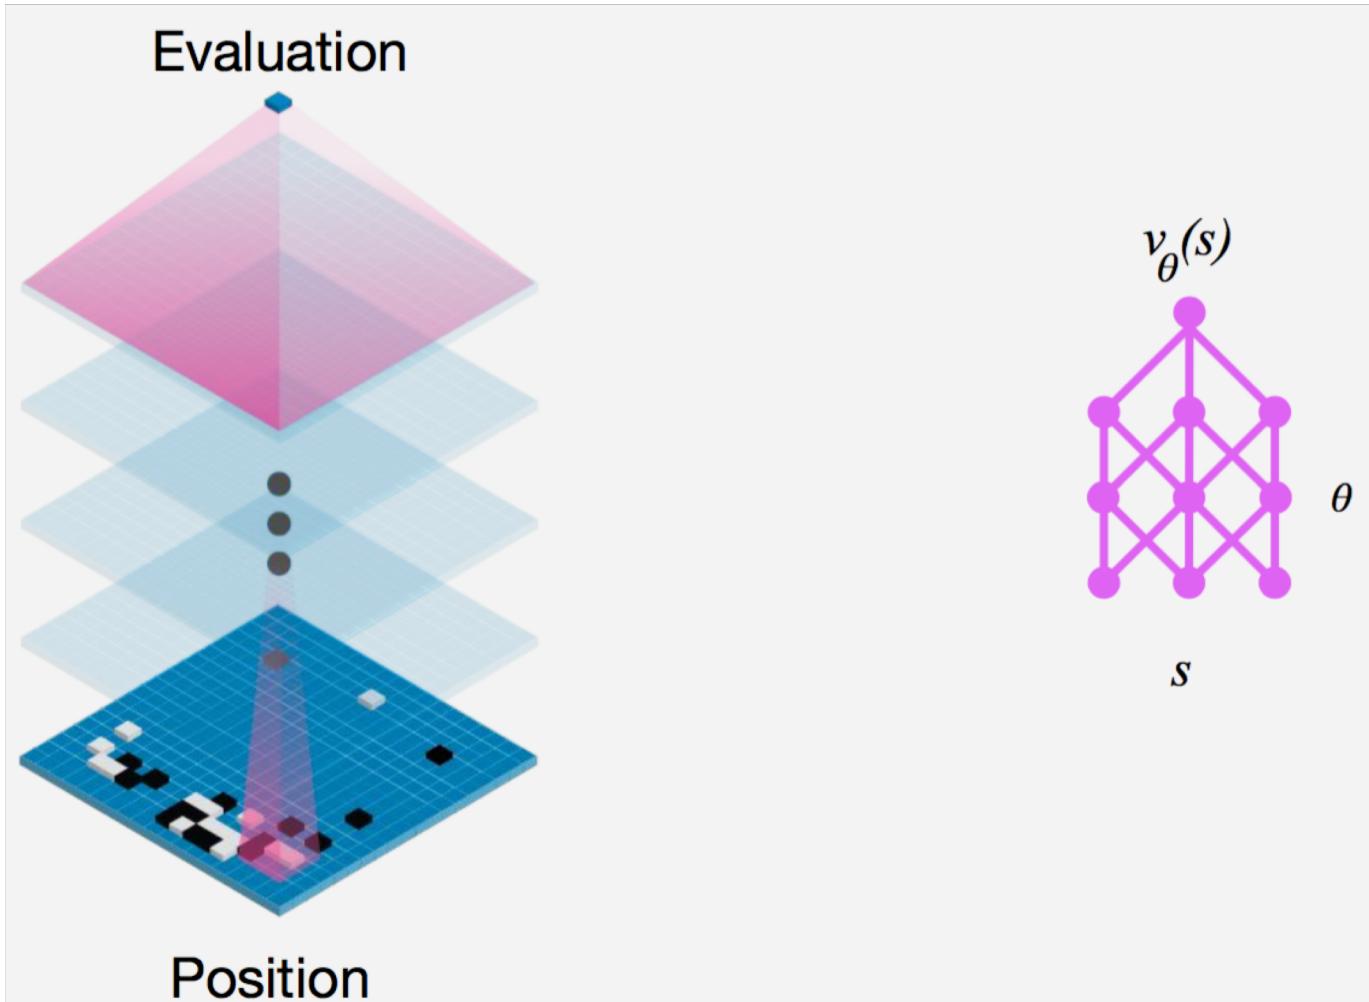
# AlphaGo Extensions

- Uses Monte Carlo tree search for action selection
- But uses a deep policy network and a deep value network to **truncate** the search tree



# AlphaGo Extensions

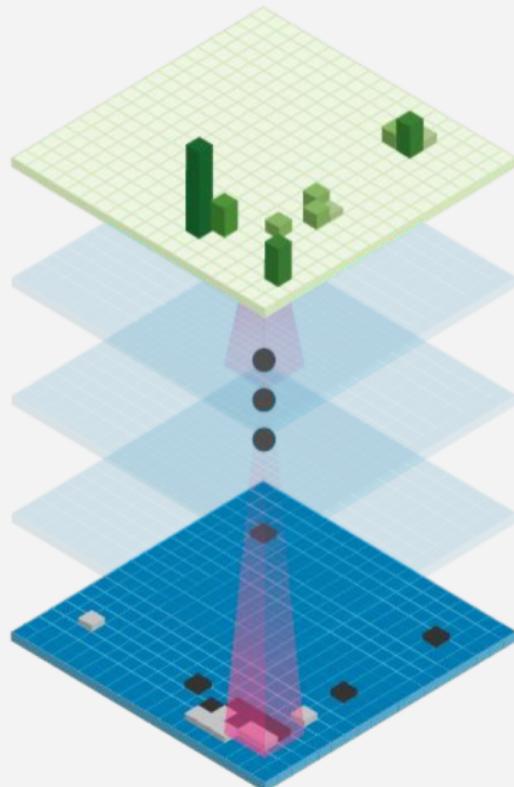
- Value Network



# AlphaGo Extensions

- Policy Network

Move probabilities

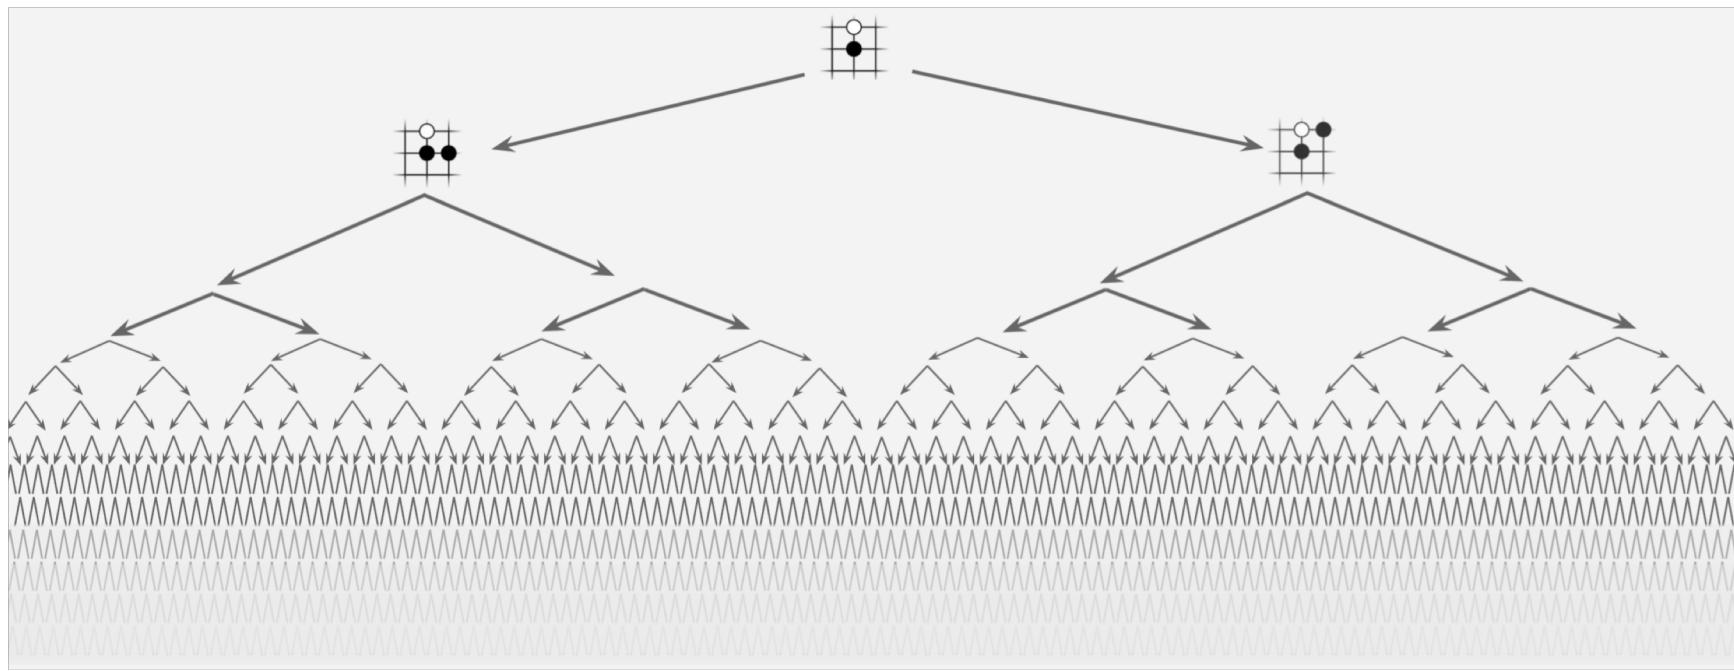


$$p_{\sigma}(a|s)$$

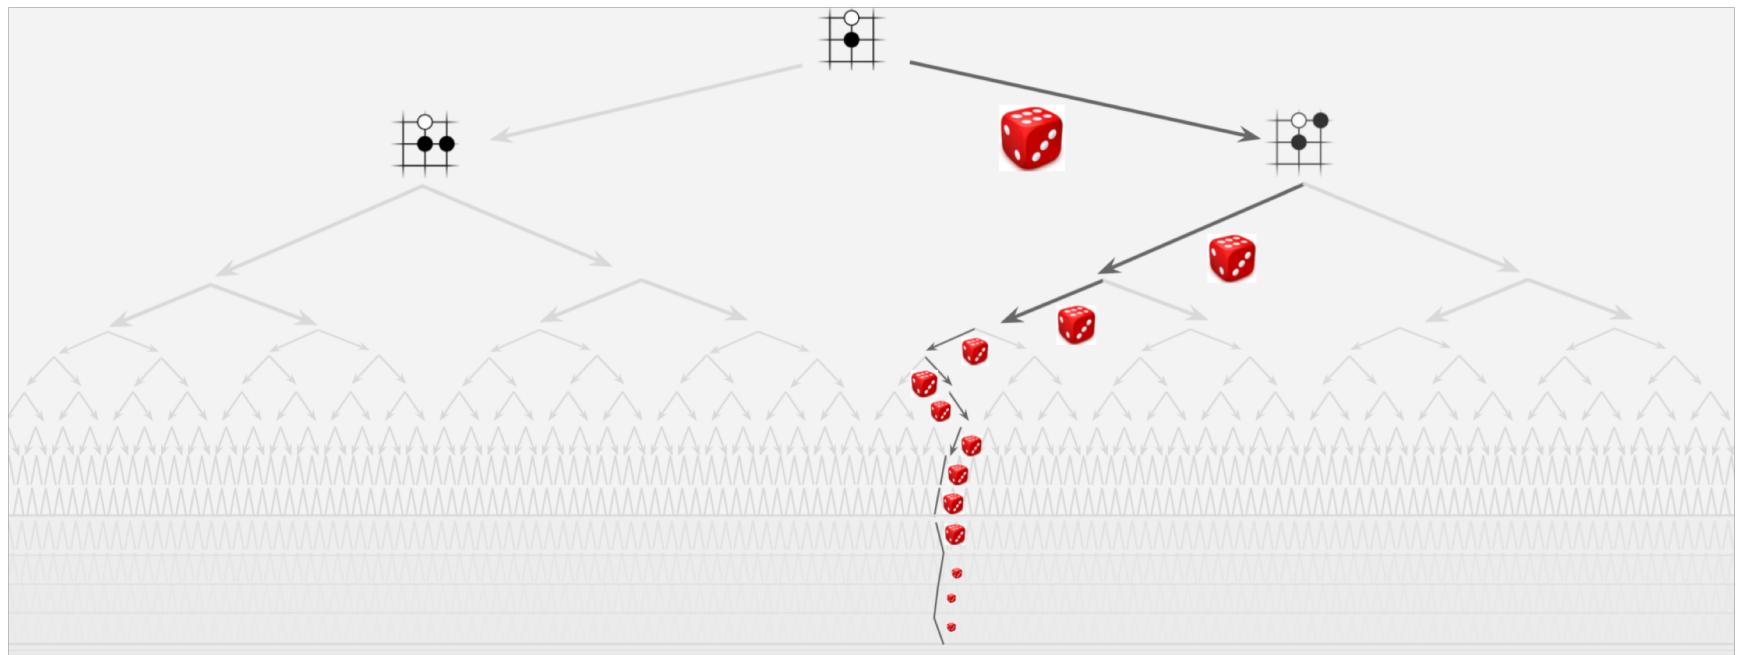
$s$

$\sigma$

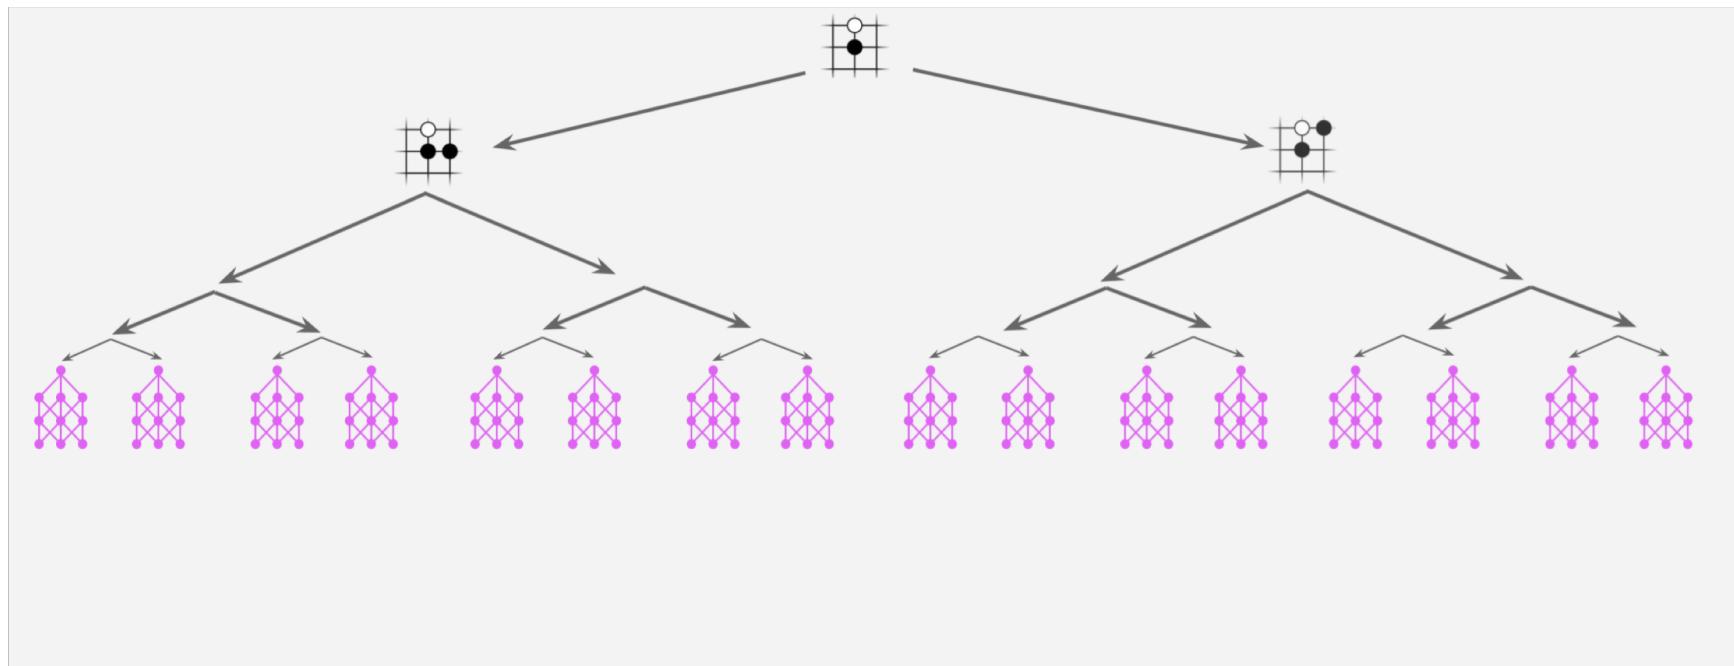
# AlphaGo Extensions



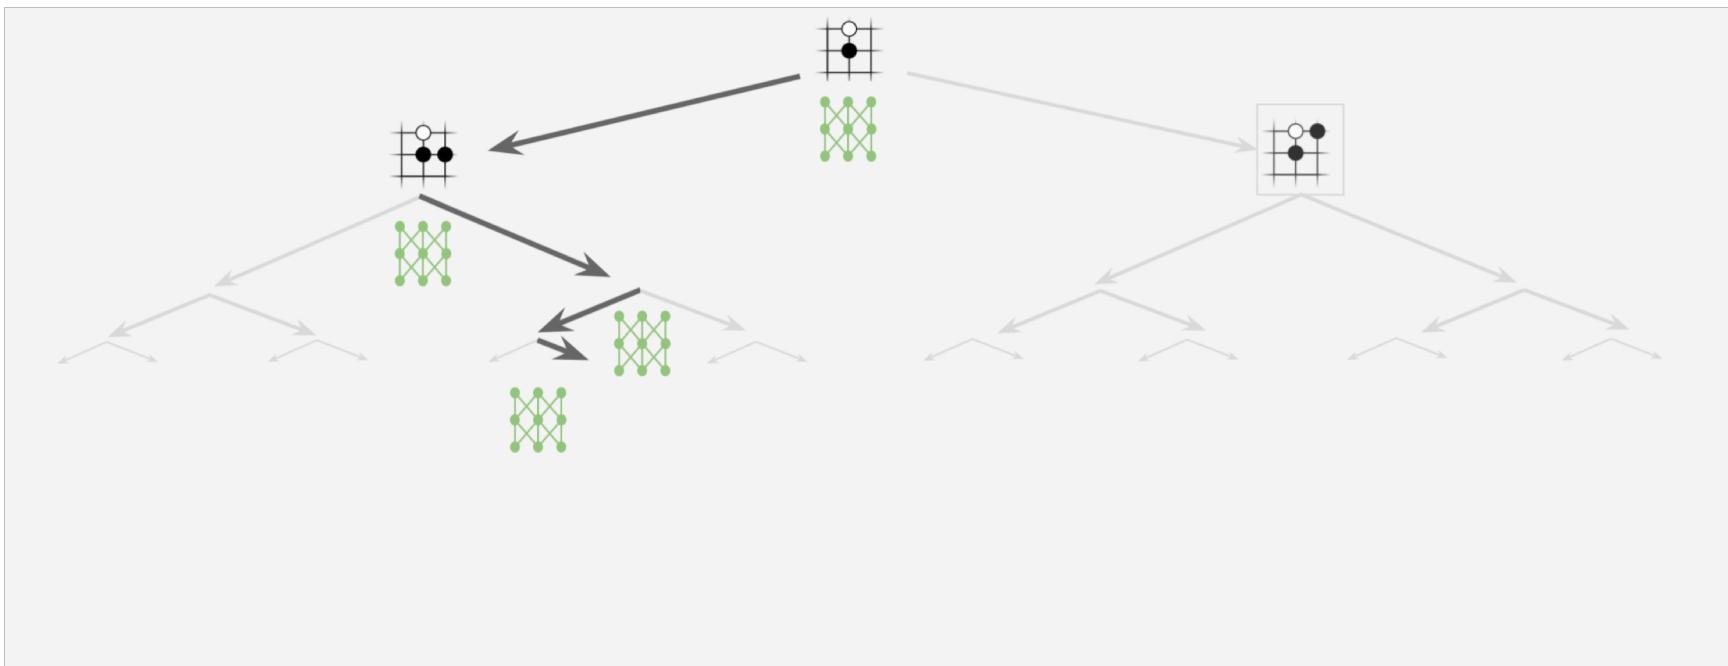
# AlphaGo Extensions



# AlphaGo Extensions

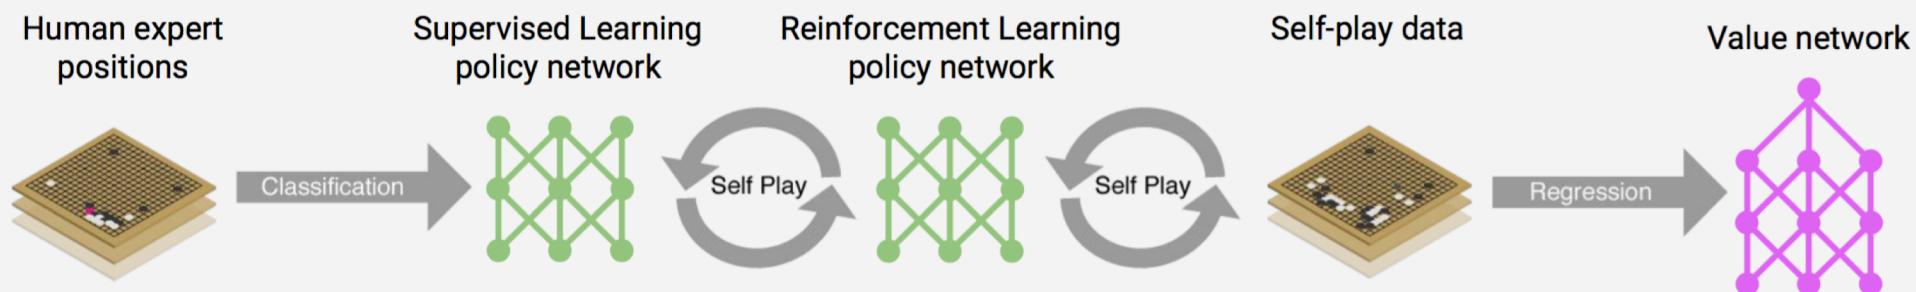


# AlphaGo Extensions



# AlphaGo Extensions

- Training the two networks



# AlphaGo Extensions

- The initial policy network

**Policy network:** 12 layer convolutional neural network

**Training data:** 30M positions from human expert games (KGS 5+ dan)

**Training algorithm:** maximise likelihood by stochastic gradient descent

$$\Delta\sigma \propto \frac{\partial \log p_\sigma(a|s)}{\partial \sigma}$$



**Training time:** 4 weeks on 50 GPUs using Google Cloud

**Results:** 57% accuracy on held out test data (state-of-the art was 44%)

# AlphaGo Extensions

- The final policy network

**Policy network:** 12 layer convolutional neural network

**Training data:** games of self-play between policy network

**Training algorithm:** maximise wins  $z$  by policy gradient reinforcement learning

$$\Delta\sigma \propto \frac{\partial \log p_\sigma(a|s)}{\partial \sigma} z$$



**Training time:** 1 week on 50 GPUs using Google Cloud

**Results:** 80% vs supervised learning. Raw network ~3 amateur dan.

# AlphaGo Extensions

- The value network

**Value network:** 12 layer convolutional neural network

**Training data:** 30 million games of self-play

**Training algorithm:** minimise MSE by stochastic gradient descent

$$\Delta\theta \propto \frac{\partial v_\theta(s)}{\partial \theta} (z - v_\theta(s))$$

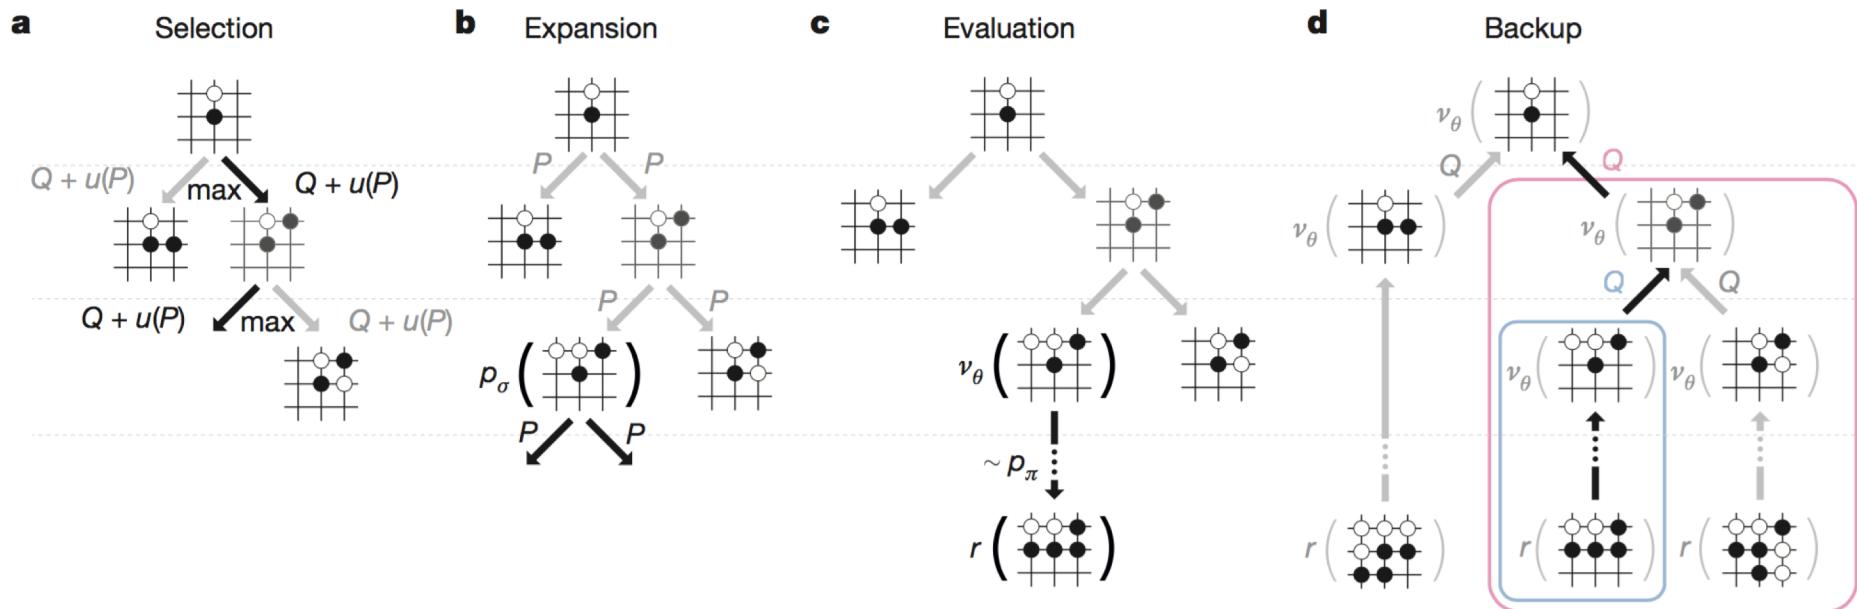


**Training time:** 1 week on 50 GPUs using Google Cloud

**Results:** First strong position evaluation function - previously thought impossible

# AlphaGo Extensions

- The MCTS procedure



# AlphaGo



**AlphaGO**

1202 CPUs, 176 GPUs,  
100+ Scientists.

**Lee Se-dol**

1 Human Brain,  
1 Coffee.

---

---

# Questions?

# Summary

---

- RL is a great framework to make agents intelligent
  - Specify goals and provide feedback
  - Traditional methods are not scalable
- Function approximation lets us manage scale (number of states)
- Complements deep learning (that solves the perception problem) allowing practical AI agents
  - DQN: Experience replay, freezed Q-targets
  - AlphaGo: Monte Carlo Tree Search with approximations
- Many challenges still remain
  - Inefficient exploration, partial observability etc.



# Appendix

# Sample Exam Questions

---

- What is the purpose of function approximation?
- Can state value function be function approximated? Is the data in the replay memory i.i.d.?
- What is a search tree? Why is it used?
- How are simulations used in a forward search? (i.e., in a simple Monte Carlo search)
- What are some practical issues with deploying an RL agent in real world?

# Additional Resources

---

- An Introduction to Reinforcement Learning by Richard Sutton and Andrew Barto
  - <http://incompleteideas.net/sutton/book/the-book.html>
- Course on Reinforcement Learning by David Silver at UCL (includes video lectures)
  - <http://www.cs.ucl.ac.uk/staff/d.silver/web/Teaching.html>
- Research Papers
  - Deep RL collection: <https://github.com/junhyukoh/deep-reinforcement-learning-papers>
  - [MKSRVBGRFOPBSAKKWLH2015] Mnih et al. Human-level control through deep reinforcement learning. *Nature*, 518:529–533, 2015.
  - [SHMGSDSAPLDGNKSLLKGH2016] Silver et al. Mastering the game of Go with deep neural networks and tree search. *Nature*, 529: 484–489, 2016.

# Recap of DQN Extensions

---

- Experience replay
  - Store transitions in replay memory  $D$
  - Sample a subset from  $D$
  - Optimize mean squared error between
    - $R_{t+1} + \gamma \max_{a \in A} Q(S_{t+1}, a, w)$  and  $Q(S_t, A_t, w)$  on this data
- Fixed Q-targets
  - Fix parameter  $w$  in  $R_{t+1} + \gamma \max_{a \in A} Q(S_{t+1}, a, w)$  for several steps

# Cons of RL

---

- In general, Reinforcement Learning requires experiencing the environment many many times
  - This is because it is a trial and error based approach
- 
- May be impractical for many complex tasks
  - Unless one has access to simulators where an RL agent can practice a billion+ times

# RL Topics Not Covered

---

- Partial observability of states
- Monte Carlo methods
  - Example:  $\epsilon$ -Greedy Policy Iteration with Monte Carlo estimation
- Temporal difference methods
  - Example: SARSA( $\lambda$ )
- Policy function approximation
- Model based methods
- ...